

## 1. Importing necessary libraries

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
In [3]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [4]: from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, train_te
```

## 2. DATA Loading

```
In [5]: df_m = pd.read_excel("premiums.xlsx")
df_m.head(10)
```

Out[5]:

	<b>Age</b>	<b>Gender</b>	<b>Region</b>	<b>Marital_Status</b>	<b>Number Of Dependents</b>	<b>BMI_Category</b>	<b>Smoking_Status</b>
<b>0</b>	26	Male	Northwest	Unmarried	0	Normal	No Smoking
<b>1</b>	29	Female	Southeast	Married	2	Obesity	Occasional
<b>2</b>	49	Female	Northeast	Married	2	Normal	No Smoking
<b>3</b>	30	Female	Southeast	Married	3	Normal	No Smoking
<b>4</b>	18	Male	Northeast	Unmarried	0	Overweight	Occasional
<b>5</b>	56	Male	Northeast	Married	3	Obesity	Occasional
<b>6</b>	33	Male	Southeast	Married	3	Normal	No Smoking
<b>7</b>	43	Male	Northeast	Married	3	Overweight	Occasional
<b>8</b>	59	Female	Southeast	Unmarried	0	Overweight	No Smoking
<b>9</b>	22	Female	Northwest	Unmarried	0	Underweight	No Smoking

In [6]: `df_m.shape`

Out[6]: (50000, 13)

In [7]: `df_m.columns`

Out[7]: Index(['Age', 'Gender', 'Region', 'Marital\_Status', 'Number Of Dependents', 'BMI\_Category', 'Smoking\_Status', 'Employment\_Status', 'Income\_Level', 'Income\_Lakhs', 'Medical History', 'Insurance\_Plan', 'Annual\_Premium\_Amount'],  
dtype='object')

In [8]: *# giving the column names a more proper structure*

```
df_m.columns = df_m.columns.str.replace(" ", "_").str.lower()
df_m.head(5)
```

```
Out[8]:   age  gender    region  marital_status  number_of_dependants  bmi_category
0     26    Male  Northwest      Unmarried                  0        Normal
1     29  Female  Southeast      Married                   2       Obesity
2     49  Female  Northeast      Married                   2        Normal
3     30  Female  Southeast      Married                   3        Normal
4     18    Male  Northeast      Unmarried                  0    Overweight
```

### 3. Exploratory Data Analysis and Data Cleaning

#### 3A - Handling the NA Values

```
In [9]: df_m.isna().sum()
```

```
Out[9]: age          0
         gender        0
         region        0
         marital_status 0
         number_of_dependants 0
         bmi_category   0
         smoking_status 11
         employment_status 2
         income_level    13
         income_lakhs     0
         medical_history  0
         insurance_plan  0
         annual_premium_amount 0
         dtype: int64
```

```
In [10]: df_m.dropna(inplace=True)
```

```
In [11]: df_m.shape
```

```
Out[11]: (49976, 13)
```

```
In [12]: df_m.isna().sum()
```

```
Out[12]: age          0  
         gender        0  
         region        0  
         marital_status 0  
         number_of_dependants 0  
         bmi_category    0  
         smoking_status   0  
         employment_status 0  
         income_level      0  
         income_lakhs       0  
         medical_history    0  
         insurance_plan     0  
         annual_premium_amount 0  
         dtype: int64
```

### 3B - Handling the Duplicates

```
df_m.duplicated().sum()
```

```
In [13]: df_m.drop_duplicates(inplace=True)  
df_m.duplicated().sum()
```

```
Out[13]: 0
```

```
In [14]: df_m.describe()
```

```
Out[14]:      age  number_of_dependants  income_lakhs  annual_premium_am  
count  49976.000000           49976.000000  49976.000000        49976.00  
mean   34.591764            1.711842   23.021150      15766.81  
std    15.000378            1.498195   24.221794      8419.99  
min    18.000000           -3.000000   1.000000      3501.00  
25%    22.000000           0.000000   7.000000      8607.75  
50%    31.000000           2.000000  17.000000     13928.00  
75%    45.000000           3.000000  31.000000     22273.50  
max    356.000000          5.000000  930.000000     43471.00
```

**number\_of\_dependents**, the min. value is -3 (**wrong**)

### 3C - Data Cleaning

```
In [15]: df_m[df_m['number_of_dependants']<0]['number_of_dependants'].unique()
```

```
Out[15]: array([-3, -1], dtype=int64)
```

We can see some negative values in number\_of\_dependants. We can replace them with positive numbers

```
In [16]: df_m['number_of_dependants'] = df_m['number_of_dependants'].abs()

df_m.describe()
```

```
Out[16]:
```

	age	number_of_dependants	income_lakhs	annual_premium_am
<b>count</b>	49976.000000	49976.000000	49976.000000	49976.00
<b>mean</b>	34.591764	1.717284	23.021150	15766.81
<b>std</b>	15.000378	1.491953	24.221794	8419.99
<b>min</b>	18.000000	0.000000	1.000000	3501.00
<b>25%</b>	22.000000	0.000000	7.000000	8607.75
<b>50%</b>	31.000000	2.000000	17.000000	13928.00
<b>75%</b>	45.000000	3.000000	31.000000	22273.50
<b>max</b>	356.000000	5.000000	930.000000	43471.00

## 3D - Numeric Columns

### 3D(i) - Univariate Analysis : Numeric Column

#### WHY ??

I use **select\_dtypes** to dynamically identify numeric features, which helps streamline preprocessing steps like scaling, correlation analysis, and feature selection in regression models.

```
In [17]: numeric_columns = df_m.select_dtypes(include=['float64', 'int64']).columns
numeric_columns

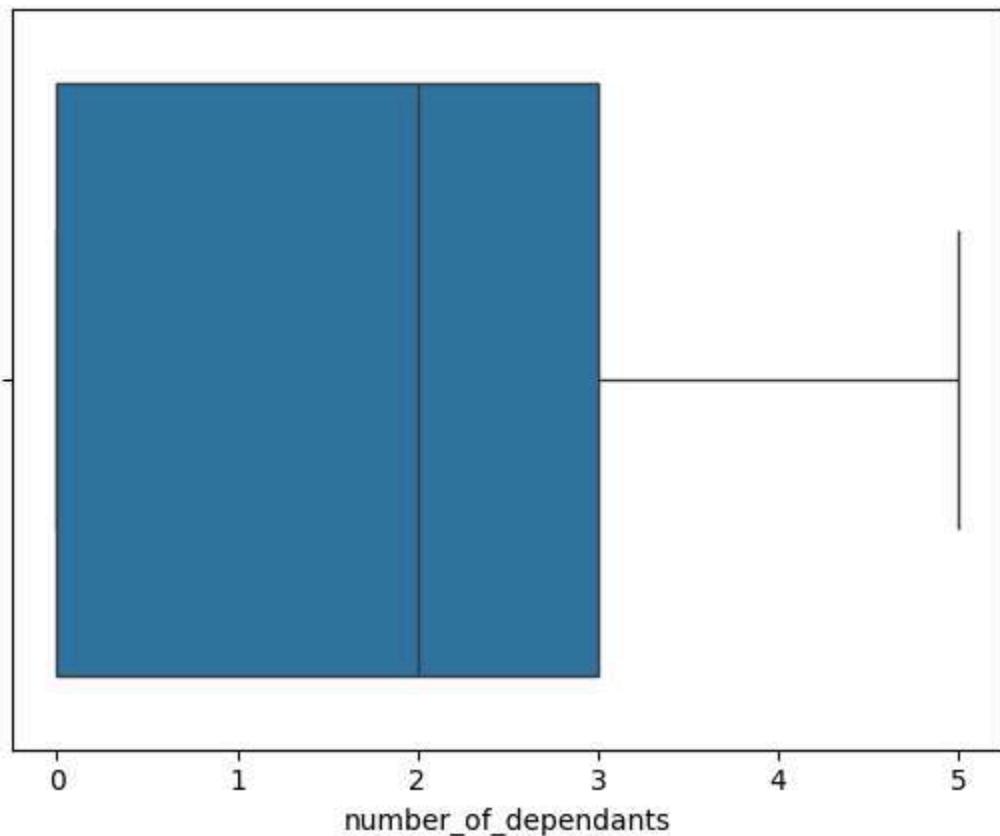
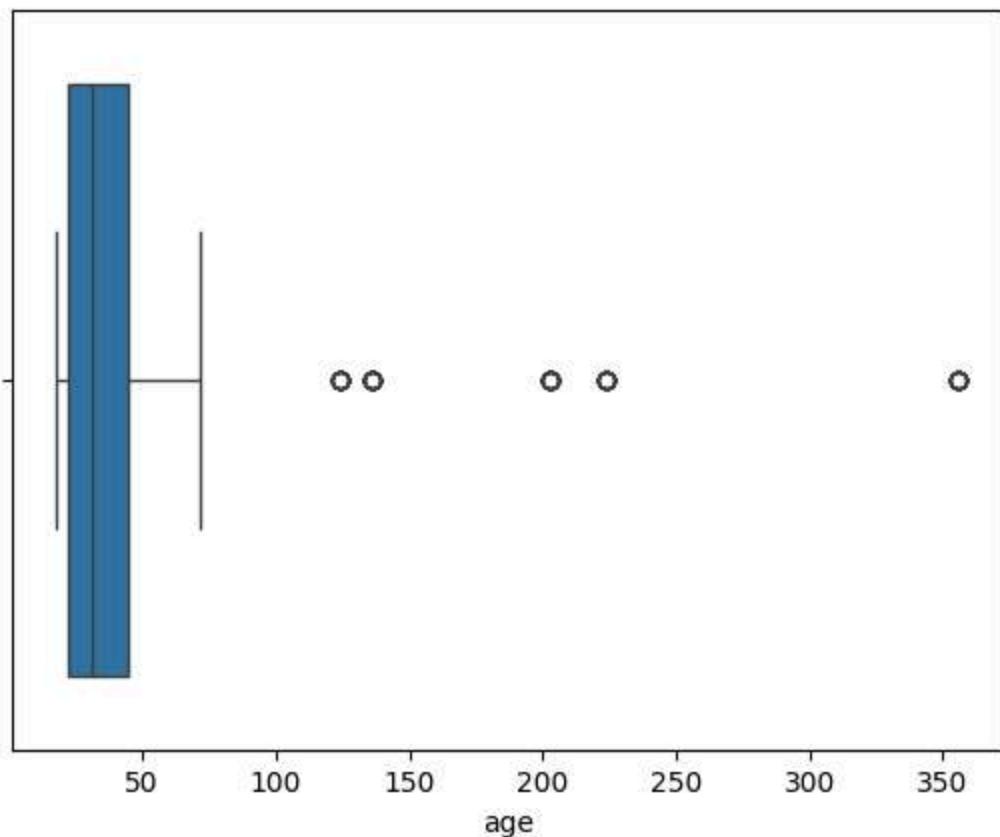
# It automatically finds all numeric columns (integers + decimals) from df_m a
```

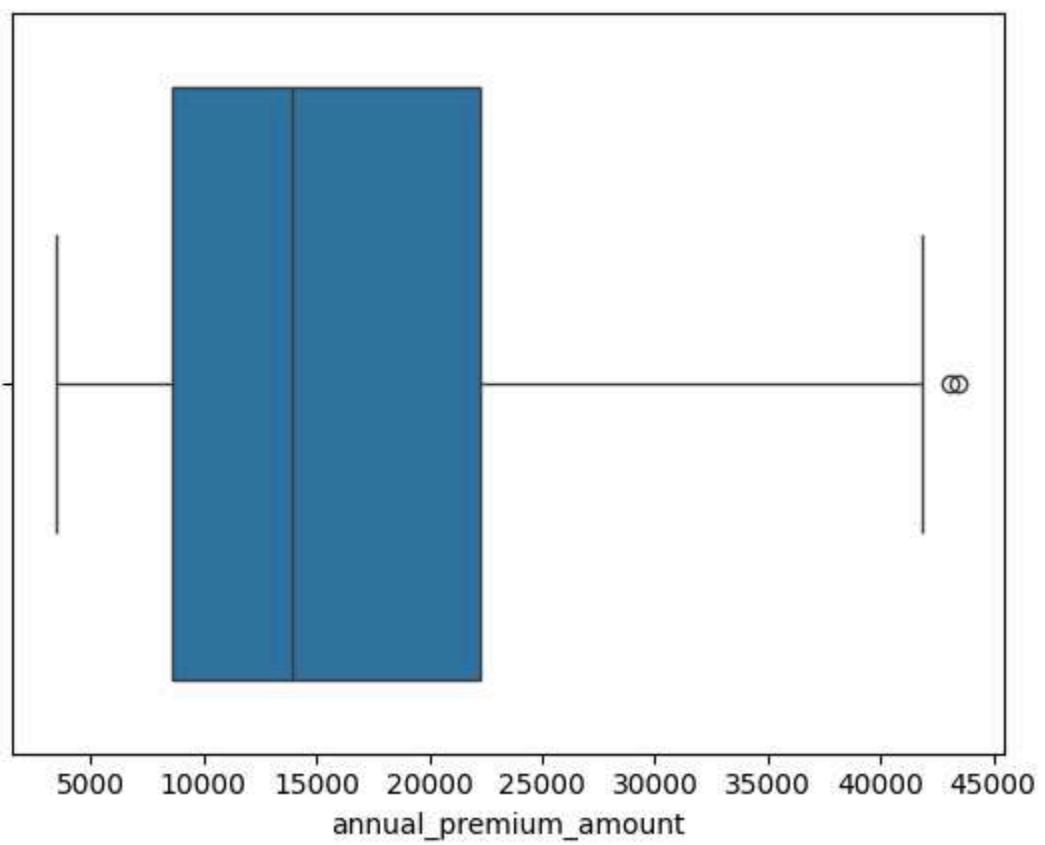
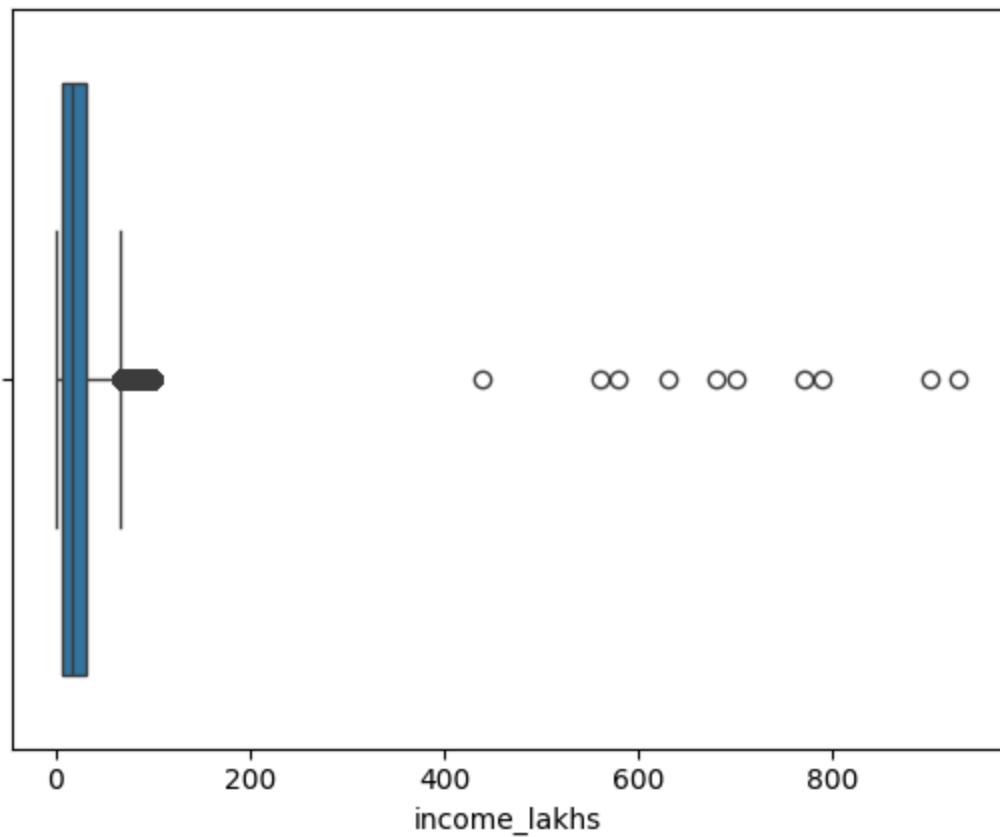
```
Out[17]: Index(['age', 'number_of_dependants', 'income_lakhs', 'annual_premium_amoun
t'], dtype='object')
```

### 3D(ii) - Box plots for numeric columns

```
In [18]: for col in numeric_columns:
    sns.boxplot(x = df_m[col])
```

```
plt.show()
```





### 3D(iii) - Outlier Treatment

For 'Age' column

```
In [19]: df_m[df_m['age']>100]['age'].unique()  
  
Out[19]: array([224, 124, 136, 203, 356], dtype=int64)  
  
In [20]: df1 = df_m[df_m.age <= 100]  
df1.age.describe()  
  
Out[20]: count    49918.000000  
mean        34.401839  
std         13.681600  
min         18.000000  
25%        22.000000  
50%        31.000000  
75%        45.000000  
max         72.000000  
Name: age, dtype: float64
```

For 'Income' column

```
In [21]: def get_iqr_bounds(col):  
    Q1, Q3 = col.quantile([0.25, 0.75])  
    IQR = Q3-Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
    return lower_bound, upper_bound  
  
lower, upper = get_iqr_bounds(df1['income_lakhs'])  
lower, upper  
  
# It calculates the lower and upper limits beyond which values in a column are
```

```
Out[21]: (-29.0, 67.0)
```

```
In [22]: df1[df1.income_lakhs > upper].shape
```

```
Out[22]: (3559, 13)
```

**There are many legitimate records (3559) that we will get rid of if we use IQR bounds method. Hence after discussion with business we decided to use a simple quantile bound.**

```
In [23]: quantile_thresold = df1.income_lakhs.quantile(0.999)  
quantile_thresold
```

```
# It calculates the 99.9th percentile of the income_lakhs column and stores th
```

```
Out[23]: 100.0
```

For skewed variables like income, use high-percentile capping (e.g., 99.9%) to control extreme values while preserving data volume and business realism.

```
In [24]: df1[df1.income_lakhs > quantile_thresold].shape
```

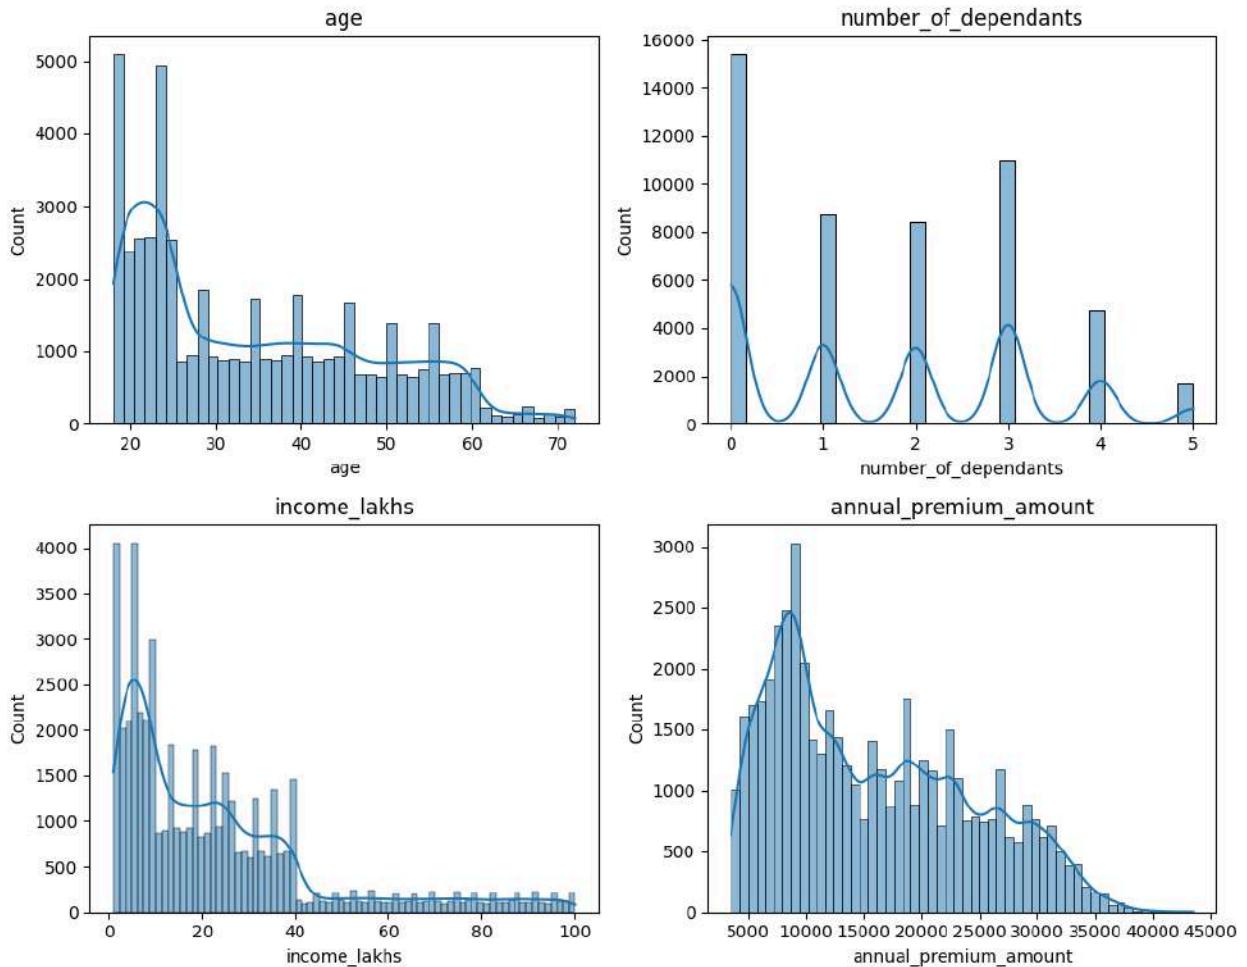
```
# It tells you how many rows have income values greater than the 99.9% quantil
```

```
Out[24]: (10, 13)
```

```
In [25]: df2 = df1[df1.income_lakhs <= quantile_thresold].copy()  
df2.describe()
```

```
Out[25]:      age  number_of_dependants  income_lakhs  annual_premium_am  
count  49908.000000                49908.000000    49908.000000          49908.00  
mean    34.401579                 1.717640     22.889897        15765.73  
std     13.681604                 1.492032     22.170699        8418.67  
min    18.000000                 0.000000     1.000000        3501.00  
25%   22.000000                 0.000000     7.000000        8608.00  
50%   31.000000                 2.000000    17.000000       13928.00  
75%   45.000000                 3.000000    31.000000       22270.50  
max   72.000000                 5.000000   100.000000      43471.00
```

```
In [26]: fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(10, 8)) # Adjust the size  
  
for i, column in enumerate(numeric_columns):  
  
    ax = axs[i//2, i%2] # Determines the position of the subplot in the grid  
    sns.histplot(df2[column], kde=True, ax=ax)  
  
    ax.set_title(column) # Each subplot gets its own column name as title  
  
plt.tight_layout() # Prevents titles and plots from overlapping  
plt.show()  
  
# The above code creates a 2x2 grid of plots and draws a histogram (with KDE)
```



### 3E - Categorical Columns

```
In [27]: categorical_cols = ['gender', 'region', 'marital_status', 'bmi_category', 'smoking_status']
for col in categorical_cols:
    print(col, ":", df2[col].unique())
```

```
gender : ['Male' 'Female']
region : ['Northwest' 'Southeast' 'Northeast' 'Southwest']
marital_status : ['Unmarried' 'Married']
bmi_category : ['Normal' 'Obesity' 'Overweight' 'Underweight']
smoking_status : ['No Smoking' 'Regular' 'Occasional' 'Smoking=0' 'Does Not Smoke'
                 'Not Smoking']
employment_status : ['Salaried' 'Self-Employed' 'Freelancer']
income_level : ['<10L' '10L - 25L' '> 40L' '25L - 40L']
medical_history : ['Diabetes' 'High blood pressure' 'No Disease'
                  'Diabetes & High blood pressure' 'Thyroid' 'Heart disease'
                  'High blood pressure & Heart disease' 'Diabetes & Thyroid'
                  'Diabetes & Heart disease']
insurance_plan : ['Bronze' 'Silver' 'Gold']
```

The category for **smoking\_status** is not uniform.

**No Smoking; Smoking=0; Does not smoke; Not Smoking** - all are SAME

```
In [28]: # giving the uniform values to the smoking_status  
  
df2['smoking_status'].replace(  
    {  
        'Not Smoking': 'No Smoking',  
        'Does Not Smoke': 'No Smoking',  
        'Smoking=0': 'No Smoking'  
    }, inplace=True)
```

```
In [29]: df2['smoking_status'].unique()
```

```
Out[29]: array(['No Smoking', 'Regular', 'Occasional'], dtype=object)
```

### 3E(i) - Univariate Analysis

```
In [30]: df2.head(5)
```

```
Out[30]:   age  gender      region  marital_status  number_of_dependants  bmi_category  
0     26    Male  Northwest      Unmarried                 0          Normal  
1     29  Female  Southeast      Married                  2          Obesity  
2     49  Female  Northeast      Married                  2          Normal  
3     30  Female  Southeast      Married                  3          Normal  
4     18    Male  Northeast      Unmarried                 0          Overweight
```

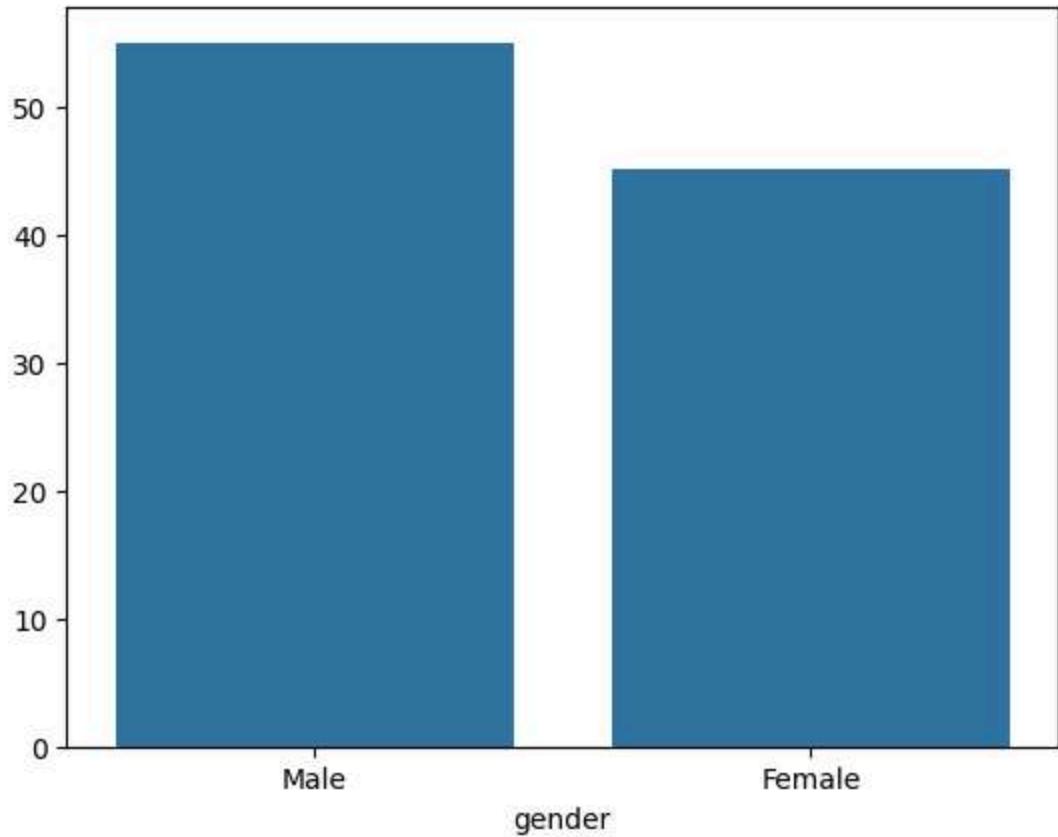
```
In [31]: # percentage count of the two genders
```

```
pct_count = df2['gender'].value_counts(normalize=True)*100  
pct_count
```

```
Out[31]: gender  
Male      54.963132  
Female    45.036868  
Name: proportion, dtype: float64
```

```
In [32]: sns.barplot(x = pct_count.index, y = pct_count.values)
```

```
Out[32]: <Axes: xlabel='gender'>
```



```
In [33]: # Checking the category-wise distribution of each category column

fig, axes = plt.subplots(3, 3, figsize=(18, 18))
axes = axes.flatten() # Flattening the 2D array of axes into 1D for easier it

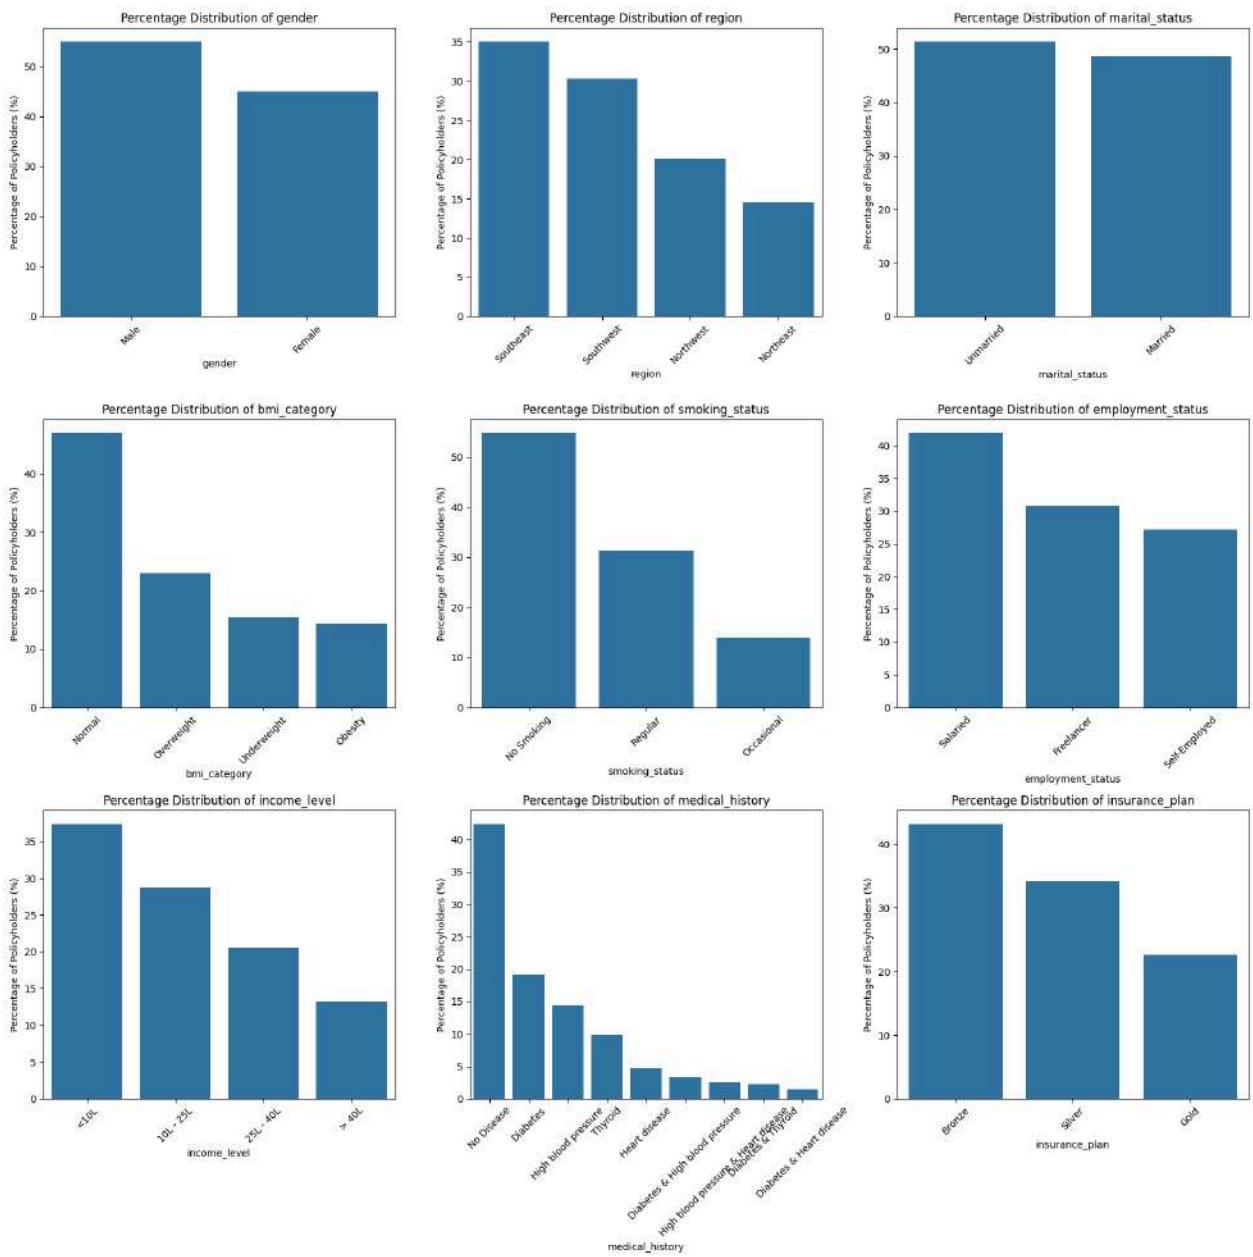
for ax, column in zip(axes, categorical_cols):

    # Calculate the percentage distribution of each category
    category_counts = df2[column].value_counts(normalize=True) * 100 # normal

    # Plotting the distribution using barplot
    sns.barplot(x = category_counts.index, y = category_counts.values, ax=ax)
    ax.set_title(f'Percentage Distribution of {column}')
    ax.set_ylabel('Percentage of Policyholders (%)')
    ax.set_xlabel(column) # Set xlabel to the column name for clarity

    ax.tick_params(axis='x', rotation=45)

plt.tight_layout() # Adjusts plot parameters for better fit in the figure window
plt.show()
```



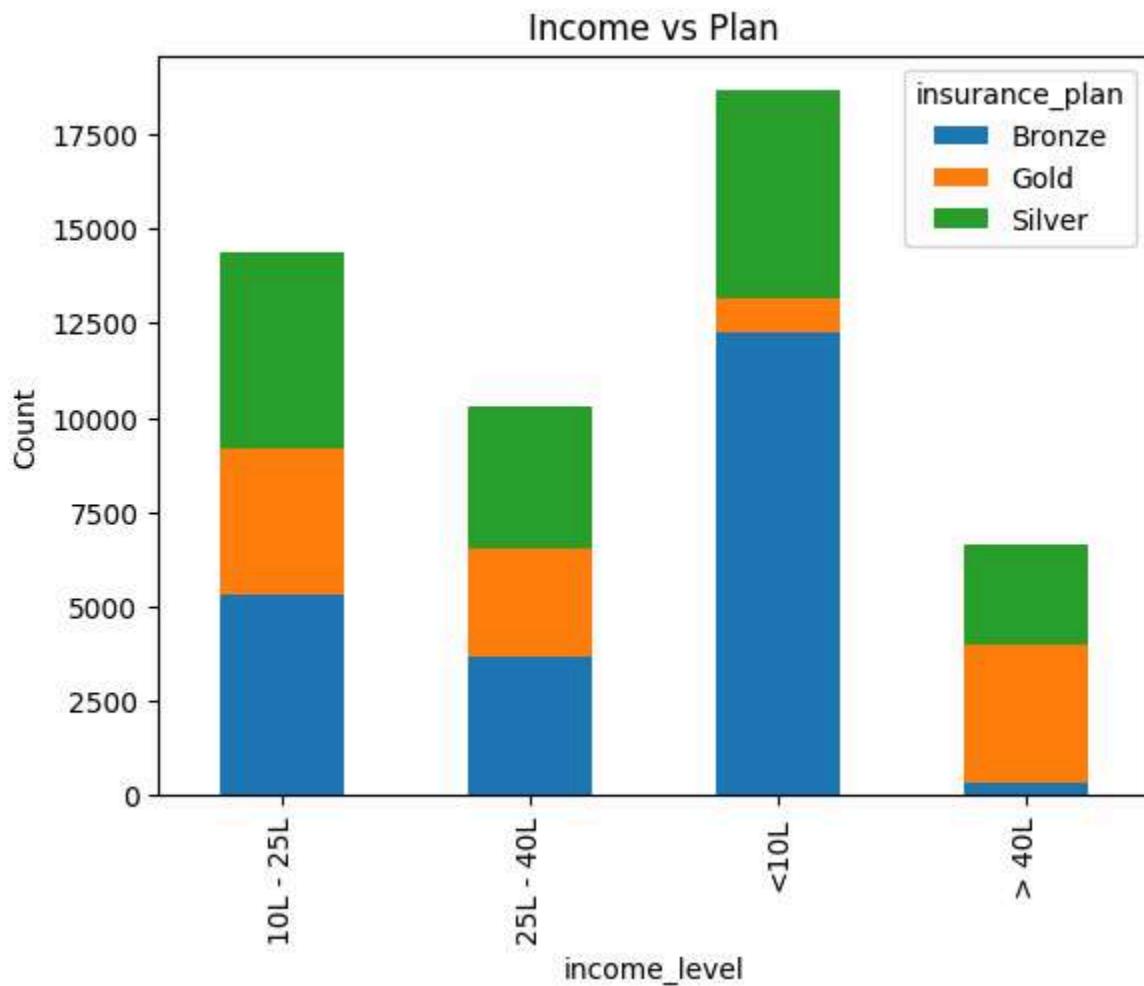
### 3E(ii) - Bivariate Analysis

In [34]: # Cross-tabulation of gender and smoking status

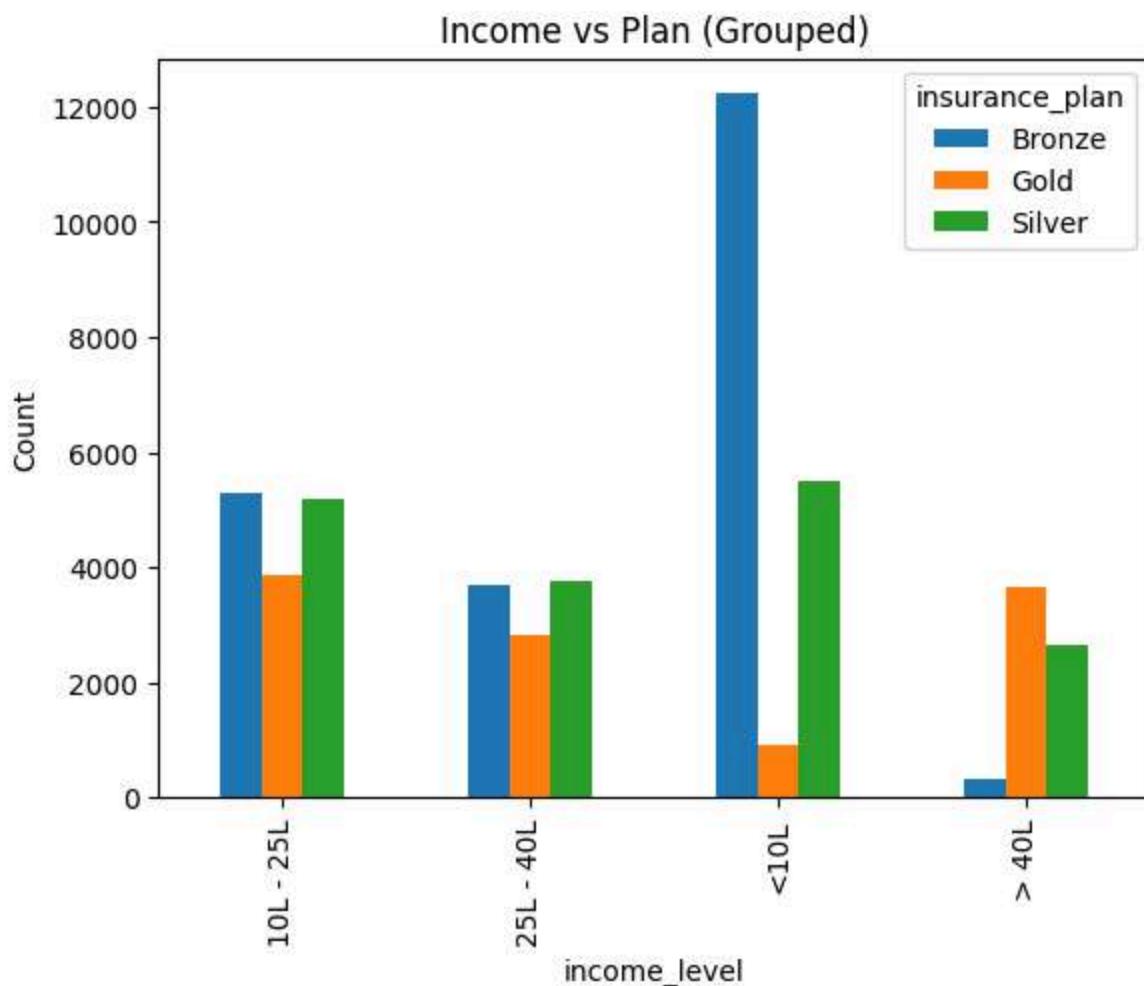
```
crosstab = pd.crosstab(df2['income_level'], df2['insurance_plan'])
print(crosstab)

# Plotting the crosstab
crosstab.plot(kind='bar', stacked=True)
plt.title('Income vs Plan')
plt.ylabel('Count')
plt.show()
```

insurance_plan	Bronze	Gold	Silver
income_level			
10L - 25L	5307	3880	5182
25L - 40L	3683	2840	3750
<10L	12226	931	5486
> 40L	329	3653	2641



```
In [35]: crosstab.plot(kind='bar')
plt.title('Income vs Plan (Grouped)')
plt.ylabel('Count')
plt.show()
```



```
In [36]: sns.heatmap(crosstab, annot=True, cmap='coolwarm', fmt="d")
plt.title('Income vs Plan')
plt.show()
```



## 4. Feature Engineering

```
In [37]: df2.head(25)
```

Out[37]:

	age	gender	region	marital_status	number_of_dependants	bmi_category
0	26	Male	Northwest	Unmarried	0	Normal
1	29	Female	Southeast	Married	2	Obesity
2	49	Female	Northeast	Married	2	Normal
3	30	Female	Southeast	Married	3	Normal
4	18	Male	Northeast	Unmarried	0	Overweight
5	56	Male	Northeast	Married	3	Obesity
6	33	Male	Southeast	Married	3	Normal
7	43	Male	Northeast	Married	3	Overweight
8	59	Female	Southeast	Unmarried	0	Overweight
9	22	Female	Northwest	Unmarried	0	Underweight
10	21	Female	Southeast	Unmarried	0	Normal
11	46	Female	Southeast	Married	4	Normal
12	68	Female	Southwest	Married	1	Normal
13	59	Female	Northeast	Married	2	Obesity
14	60	Male	Northwest	Married	5	Underweight
15	27	Male	Southwest	Married	4	Normal
16	25	Male	Southeast	Unmarried	0	Normal
17	36	Male	Northwest	Unmarried	0	Normal
18	29	Male	Northwest	Married	2	Obesity
19	20	Male	Southeast	Unmarried	2	Overweight
20	28	Female	Southwest	Unmarried	0	Overweight
21	22	Female	Southwest	Unmarried	0	Underweight
22	32	Male	Northwest	Married	4	Underweight
23	19	Male	Southwest	Unmarried	0	Normal
24	55	Male	Southeast	Married	4	Normal

In [38]: df2.medical\_history.unique()

```
Out[38]: array(['Diabetes', 'High blood pressure', 'No Disease',
       'Diabetes & High blood pressure', 'Thyroid', 'Heart disease',
       'High blood pressure & Heart disease', 'Diabetes & Thyroid',
       'Diabetes & Heart disease'], dtype=object)
```

Need to convert the above **text** data into **numeric** data

## 4A - Risk Score

It is a numerical representation of the likelihood or severity of an adverse outcome, calculated using historical data, business rules, or machine learning models.

```
In [39]: # Defining the risk scores for each condition
```

```
risk_scores = {
    "diabetes": 6,
    "heart disease": 8,
    "high blood pressure": 6,
    "thyroid": 5,
    "no disease": 0,
    "none": 0
}

df2[['disease1', 'disease2']] = df2['medical_history'].str.split(" & ", expand=True)
df2.head(10)
```

```
Out[39]:   age  gender      region  marital_status  number_of_dependants  bmi_category
  0    26    Male  Northwest        Unmarried                  0          Normal
  1    29  Female  Southeast        Married                   2          Obesity
  2    49  Female  Northeast        Married                   2          Normal
  3    30  Female  Southeast        Married                   3          Normal
  4    18    Male  Northeast        Unmarried                  0         Overweight
  5    56    Male  Northeast        Married                   3          Obesity
  6    33    Male  Southeast        Married                   3          Normal
  7    43    Male  Northeast        Married                   3         Overweight
  8    59  Female  Southeast        Unmarried                  0         Overweight
  9    22  Female  Northwest        Unmarried                  0     Underweight
```

```
In [40]: # Calculating total_risk_score

df2['disease1'].fillna('none', inplace=True)
df2['disease2'].fillna('none', inplace=True)
df2['total_risk_score'] = 0

for disease in ['disease1', 'disease2']:
    df2['total_risk_score'] += df2[disease].map(risk_scores)

max_score = df2['total_risk_score'].max()      # Normalizing the risk score to
min_score = df2['total_risk_score'].min()
df2['normalized_risk_score'] = (df2['total_risk_score'] - min_score) / (max_score - min_score)

df2.head(10)
```

	age	gender	region	marital_status	number_of_dependants	bmi_category
<b>0</b>	26	Male	Northwest	Unmarried	0	Normal
<b>1</b>	29	Female	Southeast	Married	2	Obesity
<b>2</b>	49	Female	Northeast	Married	2	Normal
<b>3</b>	30	Female	Southeast	Married	3	Normal
<b>4</b>	18	Male	Northeast	Unmarried	0	Overweight
<b>5</b>	56	Male	Northeast	Married	3	Obesity
<b>6</b>	33	Male	Southeast	Married	3	Normal
<b>7</b>	43	Male	Northeast	Married	3	Overweight
<b>8</b>	59	Female	Southeast	Unmarried	0	Overweight
<b>9</b>	22	Female	Northwest	Unmarried	0	Underweight

```
In [41]: # Encoding the other 'text' columns

df2['insurance_plan'] = df2['insurance_plan'].map({'Bronze': 1, 'Silver': 2, 'Gold': 3})
df2.head(10)
```

Out[41]:

	age	gender	region	marital_status	number_of_dependants	bmi_category
0	26	Male	Northwest	Unmarried	0	Normal
1	29	Female	Southeast	Married	2	Obesity
2	49	Female	Northeast	Married	2	Normal
3	30	Female	Southeast	Married	3	Normal
4	18	Male	Northeast	Unmarried	0	Overweight
5	56	Male	Northeast	Married	3	Obesity
6	33	Male	Southeast	Married	3	Normal
7	43	Male	Northeast	Married	3	Overweight
8	59	Female	Southeast	Unmarried	0	Overweight
9	22	Female	Northwest	Unmarried	0	Underweight

In [42]:

```
df2['income_level'] = df2['income_level'].map({'<10L':1, '10L - 25L': 2, '25L +':3})
df2.head(10)
```

Out[42]:

	age	gender	region	marital_status	number_of_dependants	bmi_category
0	26	Male	Northwest	Unmarried	0	Normal
1	29	Female	Southeast	Married	2	Obesity
2	49	Female	Northeast	Married	2	Normal
3	30	Female	Southeast	Married	3	Normal
4	18	Male	Northeast	Unmarried	0	Overweight
5	56	Male	Northeast	Married	3	Obesity
6	33	Male	Southeast	Married	3	Normal
7	43	Male	Northeast	Married	3	Overweight
8	59	Female	Southeast	Unmarried	0	Overweight
9	22	Female	Northwest	Unmarried	0	Underweight

**The other columns are 'nominal', therefore need to do One-Hot-Encoding**

In [43]: # Applying One-Hot Encoding

```
nominal_cols = ['gender', 'region', 'marital_status', 'bmi_category', 'smoking']
df3 = pd.get_dummies(df2, columns=nominal_cols, drop_first=True, dtype=int)
df3.head(3)
```

Out[43]:

	age	number_of_dependants	income_level	income_lakhs	medical_history	in:
0	26	0	1	6	Diabetes	
1	29	2	1	6	Diabetes	
2	49	2	2	20	High blood pressure	

3 rows × 23 columns

**What pd.get\_dummies() does :**

- Creates new binary columns for each category

- Replaces text with 0 or 1

In [44]: `df3.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 49908 entries, 0 to 49999
Data columns (total 23 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   age              49908 non-null   int64  
 1   number_of_dependants 49908 non-null   int64  
 2   income_level      49908 non-null   int64  
 3   income_lakhs       49908 non-null   int64  
 4   medical_history    49908 non-null   object  
 5   insurance_plan    49908 non-null   int64  
 6   annual_premium_amount 49908 non-null   int64  
 7   disease1          49908 non-null   object  
 8   disease2          49908 non-null   object  
 9   total_risk_score   49908 non-null   int64  
 10  normalized_risk_score 49908 non-null   float64 
 11  gender_Male        49908 non-null   int32  
 12  region_Northwest   49908 non-null   int32  
 13  region_Southeast    49908 non-null   int32  
 14  region_Southwest   49908 non-null   int32  
 15  marital_status_Unmarried 49908 non-null   int32  
 16  bmi_category_Obesity 49908 non-null   int32  
 17  bmi_category_Overweight 49908 non-null   int32  
 18  bmi_category_Underweight 49908 non-null   int32  
 19  smoking_status_Occasional 49908 non-null   int32  
 20  smoking_status-Regular 49908 non-null   int32  
 21  employment_status_Salaried 49908 non-null   int32  
 22  employment_status_Self-Employed 49908 non-null   int32  
dtypes: float64(1), int32(12), int64(7), object(3)
memory usage: 6.9+ MB
```

## 4B - Feature Selection

In [45]: `# Dropping the unwanted columns`

```
df4 = df3.drop(['medical_history', 'disease1', 'disease2', 'total_risk_score'],
df4.head(10)
```

Out[45]:

	age	number_of_dependants	income_level	income_lakhs	insurance_plan	annual_premium_amount
0	26	0	1	6	1	1
1	29	2	1	6	1	1
2	49	2	2	20	2	2
3	30	3	4	77	3	3
4	18	0	4	99	2	2
5	56	3	2	14	1	1
6	33	3	1	4	2	2
7	43	3	4	46	3	3
8	59	0	2	21	3	3
9	22	0	1	3	2	2

## 4C - Scaling

**I scale selected numeric features using MinMaxScaler to normalize their ranges, ensuring fair contribution to the regression model and stable optimization.**

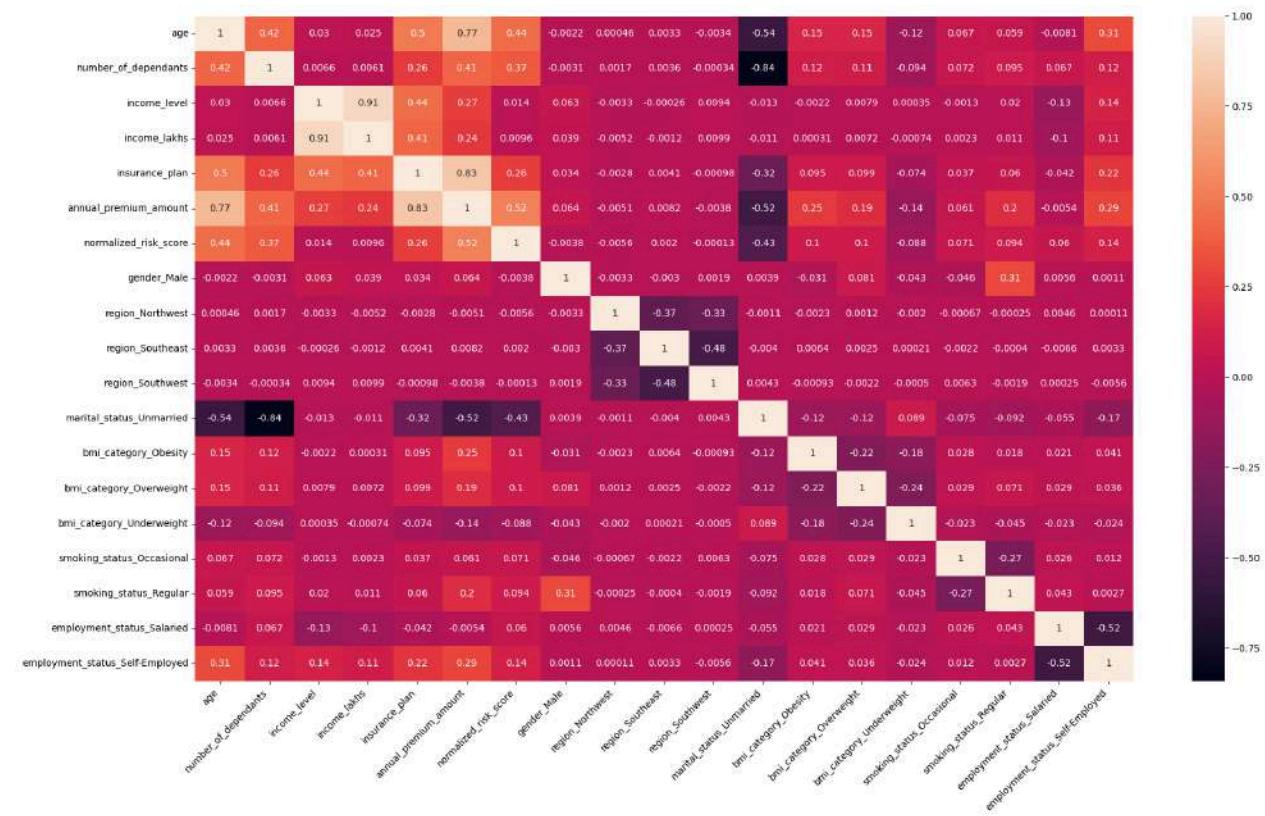
In [46]: `df4.columns`

Out[46]:

```
Index(['age', 'number_of_dependants', 'income_level', 'income_lakhs',
       'insurance_plan', 'annual_premium_amount', 'normalized_risk_score',
       'gender_Male', 'region_Northwest', 'region_Southeast',
       'region_Southwest', 'marital_status_Unmarried', 'bmi_category_Obesity',
       'bmi_category_Overweight', 'bmi_category_Underweight',
       'smoking_status_Occasional', 'smoking_status-Regular',
       'employment_status_Salaried', 'employment_status_Self-Employed'],
      dtype='object')
```

In [47]: `cm = df4.corr()`

```
plt.figure(figsize=(20,12))
sns.heatmap(cm, annot=True)
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```



```
In [48]: X = df4.drop('annual_premium_amount', axis='columns') # independent variables
y = df4['annual_premium_amount'] # target variable (what we want to predict)

from sklearn.preprocessing import MinMaxScaler # MinMaxScaler rescales the data

cols_to_scale = ['age', 'number_of_dependants', 'income_level', 'income_lakhs']
scaler = MinMaxScaler()

X[cols_to_scale] = scaler.fit_transform(X[cols_to_scale])
X.describe()
```

Out[48]:

	age	number_of_dependants	income_level	income_lakhs	insurance_plan
--	-----	----------------------	--------------	--------------	----------------

count	49908.000000	49908.000000	49908.000000	49908.000000	49908.000000
mean	0.303733		0.343528	0.365900	0.221110
std	0.253363		0.298406	0.349711	0.223946
min	0.000000		0.000000	0.000000	0.000000
25%	0.074074		0.000000	0.000000	0.060606
50%	0.240741		0.400000	0.333333	0.161616
75%	0.500000		0.600000	0.666667	0.303030
max	1.000000		1.000000	1.000000	1.000000

## 4D - Calculating VIF for Multicollinearity

```
In [49]: from statsmodels.stats.outliers_influence import variance_inflation_factor

def calculate_vif(data):
    vif_df = pd.DataFrame()
    vif_df['Column'] = data.columns
    vif_df['VIF'] = [variance_inflation_factor(data.values,i) for i in range(c
    return vif_df
```

```
In [50]: X.isna().sum().sort_values(ascending=False).head(20)
```

```
Out[50]: age                      0
number_of_dependants              0
employment_status_Salaried       0
smoking_status_Regular           0
smoking_status_Occasional        0
bmi_category_Underweight         0
bmi_category_Overweight          0
bmi_category_Obesity             0
marital_status_Unmarried         0
region_Southwest                 0
region_Southeast                 0
region_Northwest                 0
gender_Male                      0
normalized_risk_score            0
insurance_plan                   0
income_lakhs                      0
income_level                      0
employment_status_Self-Employed   0
dtype: int64
```

```
In [51]: calculate_vif(X)
```

Out[51]:

	Column	VIF
0	age	4.567634
1	number_of_dependants	4.534650
2	income_level	12.450675
3	income_lakhs	11.183367
4	insurance_plan	3.584752
5	normalized_risk_score	2.687610
6	gender_Male	2.421496
7	region_Northwest	2.102556
8	region_Southeast	2.922414
9	region_Southwest	2.670666
10	marital_status_Unmarried	3.411185
11	bmi_category_Obesity	1.352806
12	bmi_category_Overweight	1.549922
13	bmi_category_Underweight	1.302886
14	smoking_status_Occasional	1.272745
15	smoking_status-Regular	1.777089
16	employment_status_Salaried	2.382134
17	employment_status_Self-Employed	2.137753

VIF for **income\_level** and **income\_lakhs** are pretty high.

Since VIF should be <10, therefore we need to drop **one(or both)** to get the VIFs of the remaining columns within the range

In [52]: # Starting with 'income\_level'

```
calculate_vif(X.drop('income_level', axis="columns"))
```

Out[52] :

	Column	VIF
0	age	4.545825
1	number_of_dependants	4.526598
2	income_lakhs	2.480563
3	insurance_plan	3.445682
4	normalized_risk_score	2.687326
5	gender_Male	2.409980
6	region_Northwest	2.100789
7	region_Southeast	2.919775
8	region_Southwest	2.668314
9	marital_status_Unmarried	3.393718
10	bmi_category_Obesity	1.352748
11	bmi_category_Overweight	1.549907
12	bmi_category_Underweight	1.302636
13	smoking_status_Occasional	1.272744
14	smoking_status-Regular	1.777024
15	employment_status_Salaried	2.374628
16	employment_status_Self-Employed	2.132810

Since now the VIF is <5 for all columns, therefore, no need to drop 'income\_lakhs'

In [53]:

```
X_reduced = X.drop('income_level', axis="columns")  
X_reduced.head(10)
```

Out[53]:

	age	number_of_dependants	income_lakhs	insurance_plan	normalized_r
0	0.148148		0.0	0.050505	0.0
1	0.203704		0.4	0.050505	0.0
2	0.574074		0.4	0.191919	0.5
3	0.222222		0.6	0.767677	1.0
4	0.000000		0.0	0.989899	0.5
5	0.703704		0.6	0.131313	0.0
6	0.277778		0.6	0.030303	0.5
7	0.462963		0.6	0.454545	1.0
8	0.759259		0.0	0.202020	1.0
9	0.074074		0.0	0.020202	0.5

## 5. Model Training

In [54]:

```
X_train, X_test, y_train, y_test = train_test_split(X_reduced, y, test_size = 0.25)

# shape of the X_train, X_test, y_train, y_test features

print("x train: ",X_train.shape)
print("x test: ",X_test.shape)
print("y train: ",y_train.shape)
print("y test: ",y_test.shape)
```

x train: (34935, 17)  
x test: (14973, 17)  
y train: (34935,)  
y test: (14973,)

### 5A - Linear Regression Model

In [55]:

```
model_lr = LinearRegression()      # Initializes a Linear Regression model
model_lr.fit(X_train, y_train)    # It finds the best-fit line that minimizes error

test_score = model_lr.score(X_test, y_test)
train_score = model_lr.score(X_train, y_train)

train_score, test_score
```

Out[55]: (0.9282143576916762, 0.9280547230217837)

```
In [56]: # The below code checks how accurate your model's predictions are by comparing

y_pred = model_lr.predict(X_test)

mse_lr = mean_squared_error(y_test, y_pred)
rmse_lr = np.sqrt(mse_lr)
print("Linear Regression ==> MSE: ", mse_lr, "RMSE: ", rmse_lr)
```

```
Linear Regression ==> MSE:  5165611.913027982 RMSE:  2272.798256121291
```

```
In [57]: X_test.shape
```

```
Out[57]: (14973, 17)
```

```
In [58]: np.set_printoptions(suppress=True, precision=6)
model_lr.coef_
```

```
Out[58]: array([11160.926462, -676.443991, -514.229816, 12557.012936,
        4810.357702, 168.646662, -35.719292, 39.96513 ,
       -24.652929, -935.760611, 3387.911455, 1599.362268,
      391.171304, 735.912278, 2234.804712, 155.984674,
     415.903973])
```

**The above code tells Python how to display numbers, and then shows how much each feature affects the prediction in your Linear Regression model.**

**For ex :**

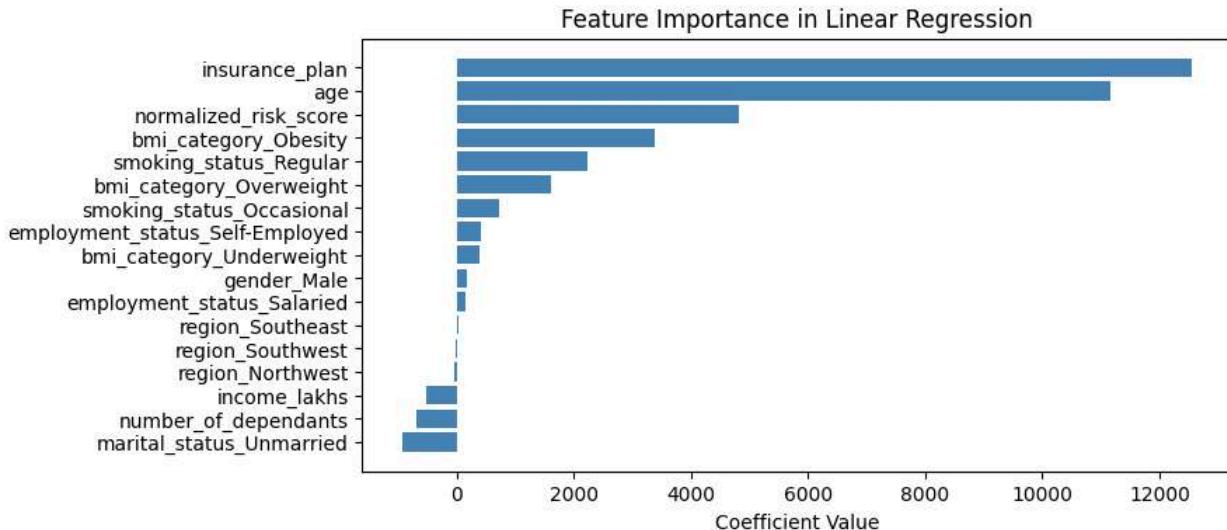
Beacuse of **suppress=True** → the output don't show scientific notation (Shows 0.000123 instead of 1.23e-4)

```
In [59]: feature_importance = model_lr.coef_

# Create a DataFrame for easier handling
coef_df = pd.DataFrame(feature_importance, index=X_train.columns, columns=['Coefficients'])

# Sort the coefficients for better visualization
coef_df = coef_df.sort_values(by='Coefficients', ascending=True)

# Plotting
plt.figure(figsize=(8, 4))
plt.barh(coef_df.index, coef_df['Coefficients'], color='steelblue')
plt.xlabel('Coefficient Value')
plt.title('Feature Importance in Linear Regression')
plt.show()
```



The above code shows which features influence the prediction the most by turning model coefficients into a clear horizontal bar chart.

Since we have around 92% of model accuracy, we need to get it higher. Therefore, we need to look into other ways to do so.

Out of which, first will be **Ridge Regression Model**.

## 5B - Ridge Regression Model

```
In [60]: model_rg = Ridge(alpha=1)

model_rg.fit(X_train, y_train)
test_score = model_rg.score(X_test, y_test)
train_score = model_rg.score(X_train, y_train)
train_score, test_score
```

Out[60]: (0.9282143198366275, 0.9280541644640345)

```
In [61]: model_rg = Ridge(alpha=10)

model_rg.fit(X_train, y_train)
test_score_rg = model_rg.score(X_test, y_test)
train_score_rg = model_rg.score(X_train, y_train)

train_score_rg, test_score_rg
```

Out[61]: (0.9282106074563636, 0.9280459054997704)

```
In [62]: model_rg = Ridge(alpha=30)

model_rg.fit(X_train, y_train)
```

```
test_score = model_rg.score(X_test, y_test)
train_score = model_rg.score(X_train, y_train)
train_score, test_score
```

```
Out[62]: (0.9281812886908575, 0.9280073438795297)
```

Not a very noticeable change, therefore, switching to another : **XG Boost**

## 5C - XG Boost

```
In [63]: from xgboost import XGBRegressor

model_xgb = XGBRegressor(n_estimators=20, max_depth=3)
model_xgb.fit(X_train, y_train)
model_xgb.score(X_test, y_test)
```

```
Out[63]: 0.9782300591468811
```

```
In [64]: y_pred = model_rg.predict(X_test)

mse_lr = mean_squared_error(y_test, y_pred)
rmse_lr = np.sqrt(mse_lr)
print("Ridge Regression ==> MSE: ", mse_lr, "RMSE: ", rmse_lr)
```

```
Ridge Regression ==> MSE:  5169013.69660484 RMSE:  2273.546501966661
```

```
In [65]: y_pred = model_xgb.predict(X_test)

mse_lr = mean_squared_error(y_test, y_pred)
rmse_lr = np.sqrt(mse_lr)
print("XGBoost Regression ==> MSE: ", mse_lr, "RMSE: ", rmse_lr)
```

```
XGBoost Regression ==> MSE:  1563064.1356043513 RMSE:  1250.2256338774819
```

The above comparison shows that : On average, **Ridge Regression's** predictions are off by about ₹2,273, but **XGBoost's** predictions are off by about ₹1,250.

That means : **XGBoost reduces error by ~₹1,000 per prediction and that's a big improvement in real terms.**

### WHY go with XG Boost ??

As we compared Ridge Regression and XGBoost using RMSE, we see that XGBoost reduced prediction error by nearly 45% compared to Ridge, indicating that the relationship between features and premium is non-linear and better captured by tree-based models.

## 5D - Setting up Randomized Search

### WHAT

I'm using **RandomizedSearchCV** to tune **XGBoost** hyperparameters by testing multiple configurations with cross-validation and selecting the model with the best R<sup>2</sup> score.

```
In [66]: model_xgb = XGBRegressor()          # creates a basic XGBoost model
param_grid = {                            # Try different combinations of these ar
    'n_estimators': [20, 40, 50],           # ...
    'learning_rate': [0.01, 0.1, 0.2],      # ...
    'max_depth': [3, 4, 5],                # ...
}
random_search = RandomizedSearchCV(model_xgb, param_grid, n_iter=10, cv=3, scoring='r2')
random_search.fit(X_train, y_train)
random_search.best_score_                 # The best average R2 score achieved during
```

```
Out[66]: 0.9809474547704061
```

```
In [67]: random_search.best_params_        # it gives the best possible combination to ha
```

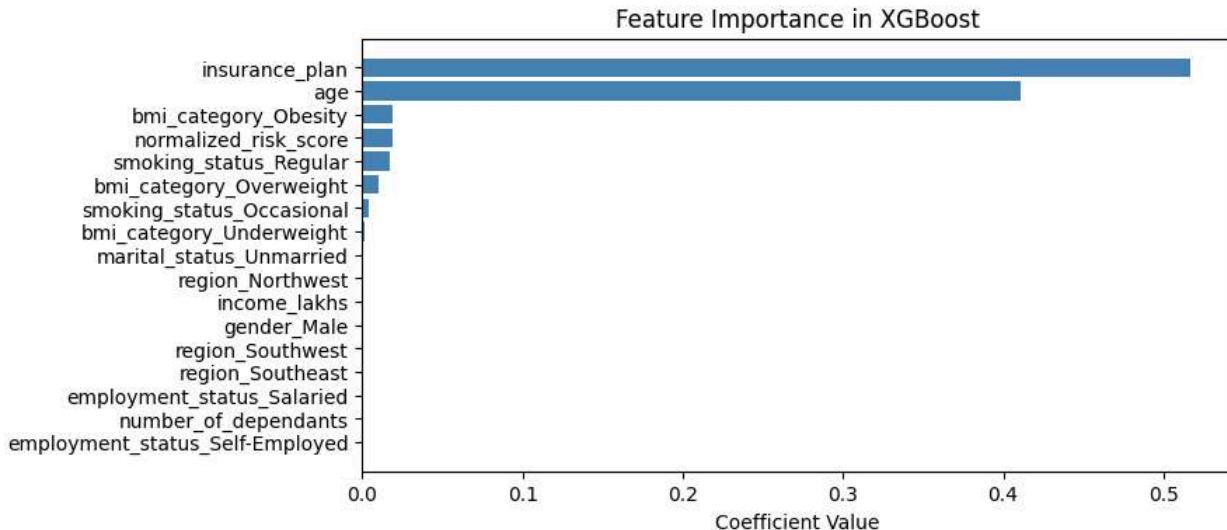
```
Out[67]: {'n_estimators': 50, 'max_depth': 5, 'learning_rate': 0.1}
```

```
In [68]: best_model = random_search.best_estimator_
```

```
In [69]: feature_importance = best_model.feature_importances_
         # Create a DataFrame for easier handling
         coef_df = pd.DataFrame(feature_importance, index=X_train.columns, columns=['Coefficients'])

         # Sort the coefficients for better visualization
         coef_df = coef_df.sort_values(by='Coefficients', ascending=True)

         # Plotting
         plt.figure(figsize=(8, 4))
         plt.barh(coef_df.index, coef_df['Coefficients'], color='steelblue')
         plt.xlabel('Coefficient Value')
         plt.title('Feature Importance in XGBoost')
         plt.show()
```



Feature Importance in **XG Boost** may not be as straightforward to interpret as Feature Importance in **Linear Regression**, since XG Boost is a **Tree-Based Model**.

## 5E - Error Analysis

```
In [70]: y_pred = best_model.predict(X_test)

residuals = y_pred - y_test
residuals_pct = (residuals / y_test) * 100

results_df = pd.DataFrame({
    'actual': y_test,
    'predicted': y_pred,
    'diff': residuals,
    'diff_pct': residuals_pct
})
results_df.head(10)
```

Out[70]:

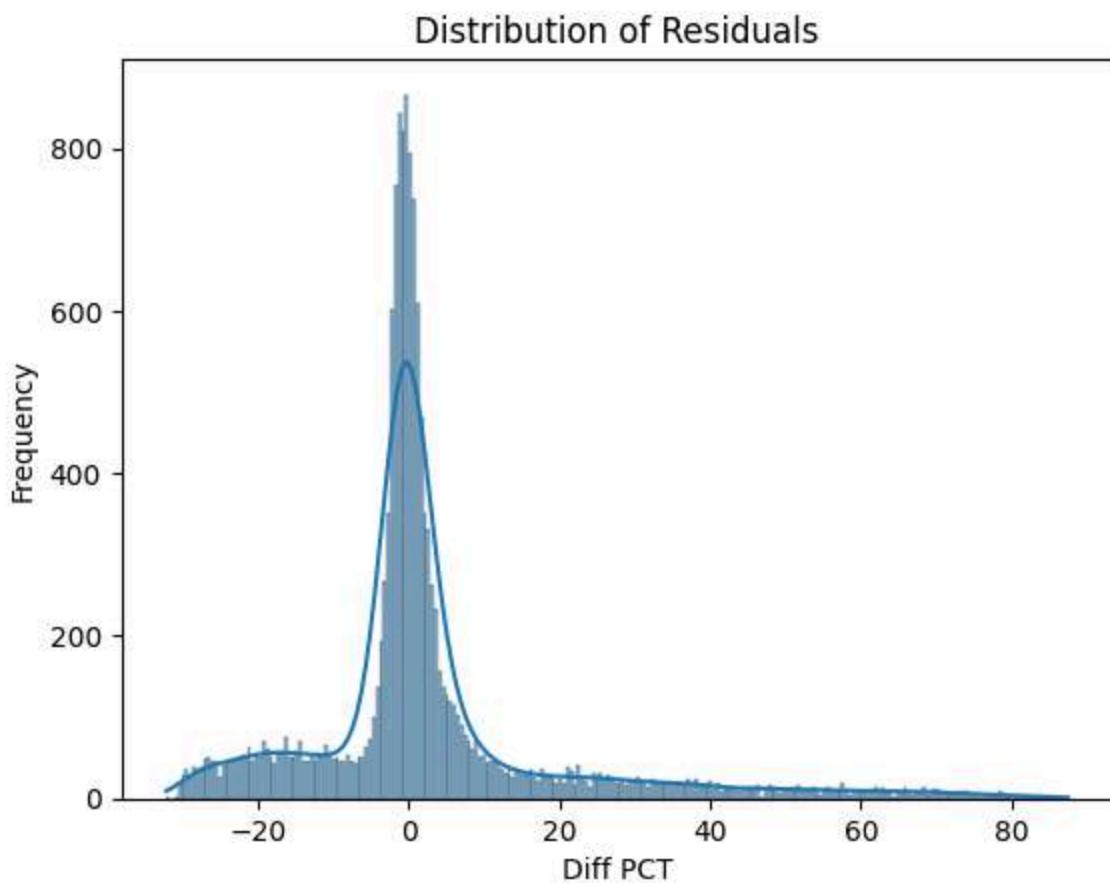
	actual	predicted	diff	diff_pct
<b>3598</b>	20554	20334.953125	-219.046875	-1.065714
<b>35794</b>	29647	29378.779297	-268.220703	-0.904714
<b>43608</b>	20560	20618.185547	58.185547	0.283004
<b>42730</b>	5018	7352.829590	2334.829590	46.529087
<b>18936</b>	8929	8203.291992	-725.708008	-8.127540
<b>45416</b>	9892	10230.427734	338.427734	3.421227
<b>20029</b>	5140	6670.849121	1530.849121	29.783057
<b>4294</b>	9631	7053.477539	-2577.522461	-26.762771
<b>39145</b>	18777	18649.619141	-127.380859	-0.678388
<b>3152</b>	14536	13996.699219	-539.300781	-3.710104

We can see above that for '42730' and '20029', the diff\_pct is higher, which is not good

In [71]:

```
# Visualizing the diff_pct

sns.histplot(results_df['diff_pct'], kde=True)
plt.title('Distribution of Residuals')
plt.xlabel('Diff PCT')
plt.ylabel('Frequency')
plt.show()
```



**There are differences that shows error upto 80%**

```
In [72]: X_test.shape
```

```
Out[72]: (14973, 17)
```

```
In [73]: # Taking out the records having 'diff_pct' >10  
extreme_error_threshold = 10  
  
extreme_results_df = results_df[np.abs(results_df['diff_pct']) > extreme_error_threshold]  
extreme_results_df.head(10)
```

Out[73]:

	<b>actual</b>	<b>predicted</b>	<b>diff</b>	<b>diff_pct</b>
<b>42730</b>	5018	7352.829590	2334.829590	46.529087
<b>20029</b>	5140	6670.849121	1530.849121	29.783057
<b>4294</b>	9631	7053.477539	-2577.522461	-26.762771
<b>44419</b>	4687	6670.849121	1983.849121	42.326629
<b>6707</b>	8826	10047.326172	1221.326172	13.837822
<b>11728</b>	4796	6565.852051	1769.852051	36.902670
<b>15740</b>	9045	10047.326172	1002.326172	11.081550
<b>35065</b>	5929	6820.886230	891.886230	15.042777
<b>9654</b>	5927	8212.278320	2285.278320	38.557083
<b>22679</b>	8482	7058.687988	-1423.312012	-16.780382

In [74]: `extreme_results_df.shape # total records having having 'diff_pct' >10`

Out[74]: (4487, 4)

In [75]: `results_df.shape`

Out[75]: (14973, 4)

In [76]: `extreme_errors_pct = extreme_results_df.shape[0]*100/X_test.shape[0]`  
`round(extreme_errors_pct, 3)`

Out[76]: 29.967

30% of the predictions are extreme error (very wrong).

That means, for 30% customers, we will either overcharge or undercharge by 10% or more.

In [77]: `extreme_results_df[abs(extreme_results_df.diff_pct)>50].sort_values("diff_pct")`

Out[77]:

	<b>actual</b>	<b>predicted</b>	<b>diff</b>	<b>diff_pct</b>
<b>36269</b>	3501	6565.852051	3064.852051	87.542189
<b>48801</b>	3516	6565.852051	3049.852051	86.742095
<b>42342</b>	3521	6565.852051	3044.852051	86.476911
<b>18564</b>	3523	6565.852051	3042.852051	86.371049
<b>7988</b>	3527	6565.852051	3038.852051	86.159684
...	...	...	...	...
<b>32671</b>	4656	6994.980957	2338.980957	50.235845
<b>14798</b>	4371	6565.852051	2194.852051	50.213957
<b>13736</b>	4371	6565.852051	2194.852051	50.213957
<b>10107</b>	4710	7073.240723	2363.240723	50.174962
<b>16908</b>	4699	7053.477539	2354.477539	50.105928

549 rows × 4 columns

There will be about 549 customers whom we will overcharge or undercharge by more than 50%

## 5F - More Analysis

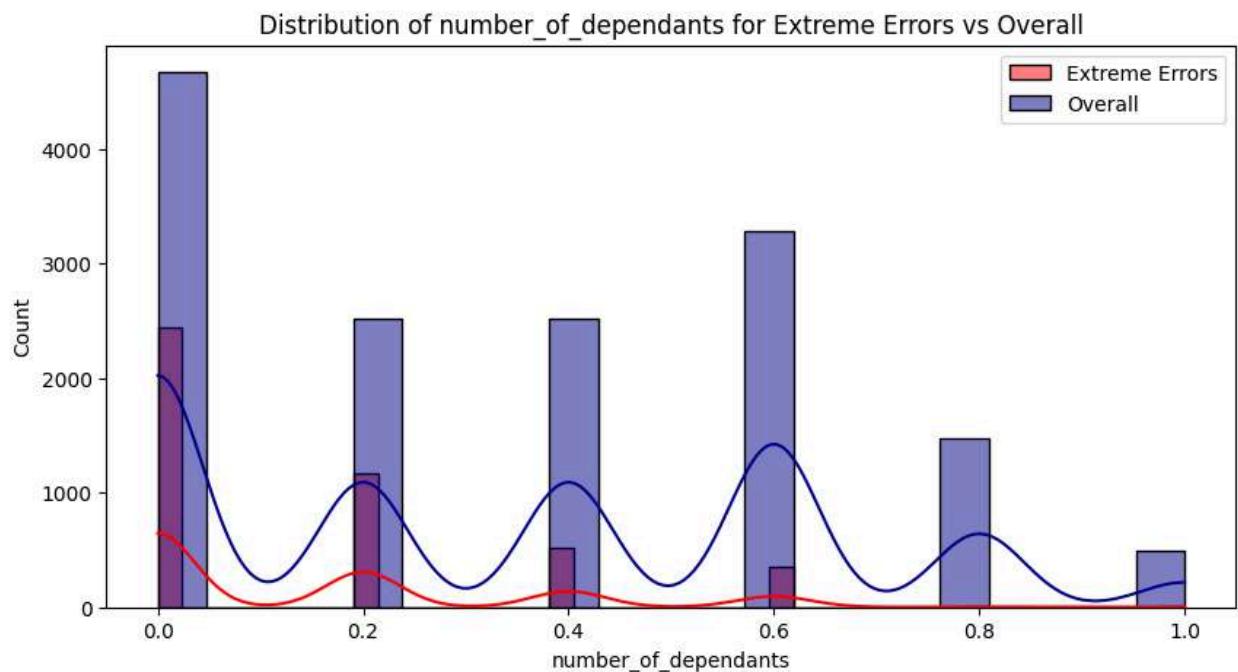
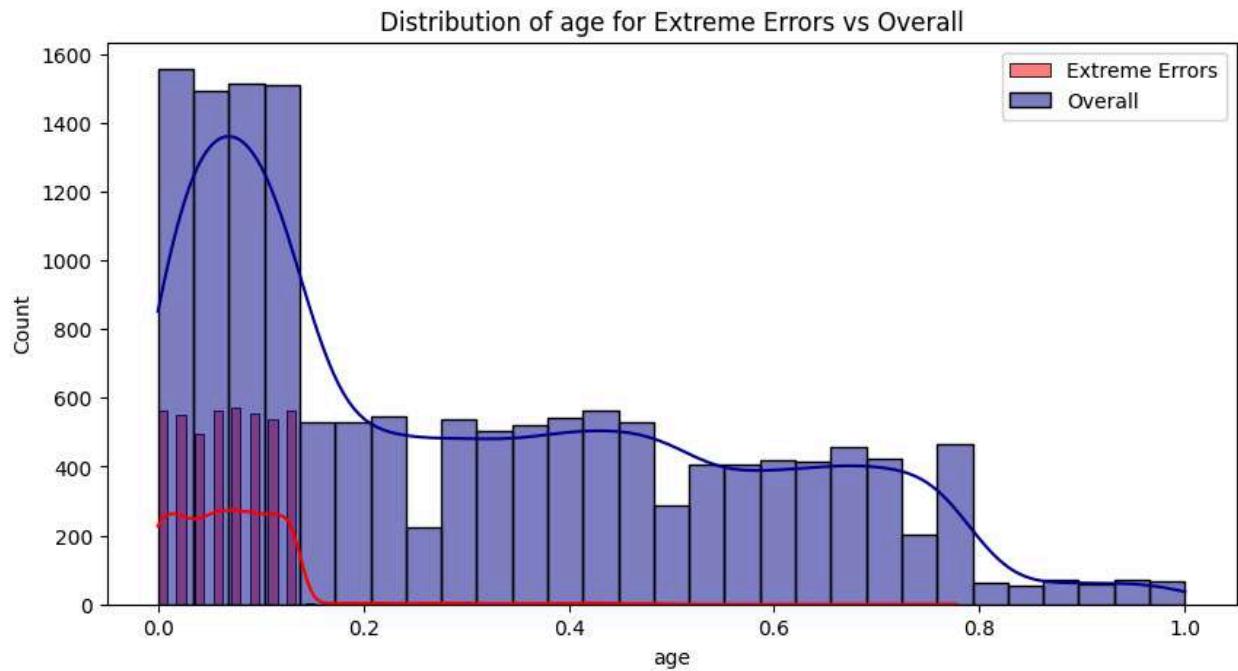
```
In [78]: extreme_errors_df = X_test.loc[extreme_results_df.index] # rows that have h  
extreme_errors_df.head(5)
```

Out[78]:

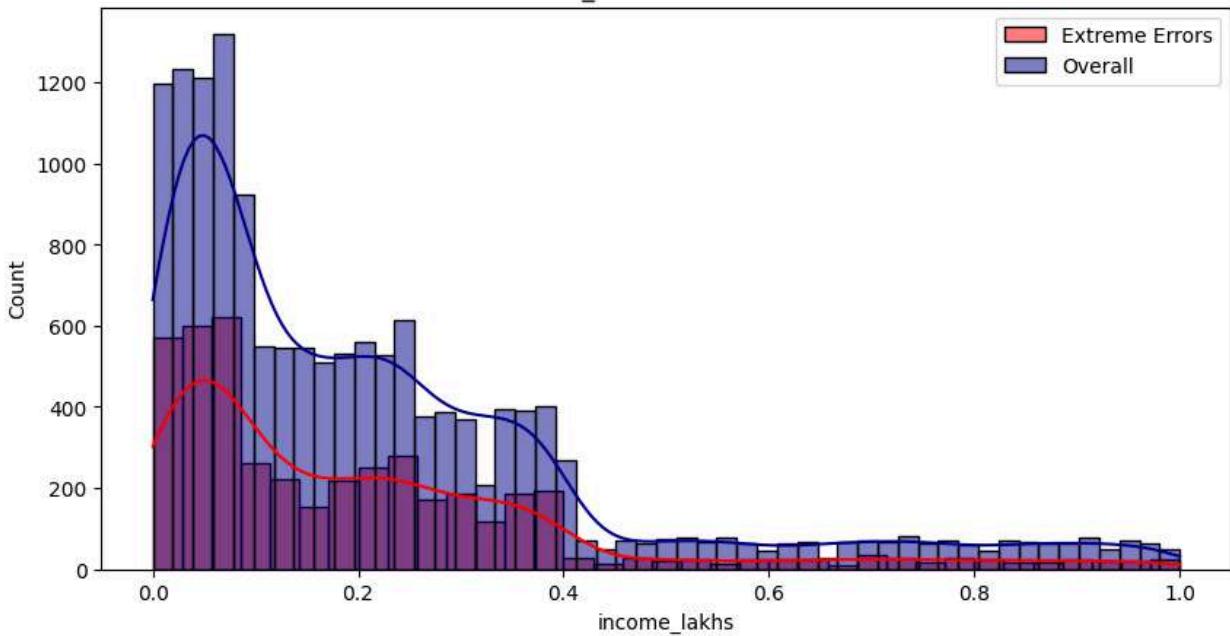
	<b>age</b>	<b>number_of_dependants</b>	<b>income_lakhs</b>	<b>insurance_plan</b>	<b>normalised</b>
<b>42730</b>	0.092593		0.2	0.131313	0.0
<b>20029</b>	0.018519		0.2	0.030303	0.0
<b>4294</b>	0.000000		0.2	0.020202	0.0
<b>44419</b>	0.055556		0.0	0.242424	0.0
<b>6707</b>	0.111111		0.2	0.070707	0.5

```
In [79]: for feature in X_test.columns:  
    plt.figure(figsize=(10, 5))  
  
    sns.histplot(extreme_errors_df[feature], color='red', label='Extreme Error')  
    sns.histplot(X_test[feature], color='darkblue', label='Overall', alpha=0.5)  
  
    plt.legend()
```

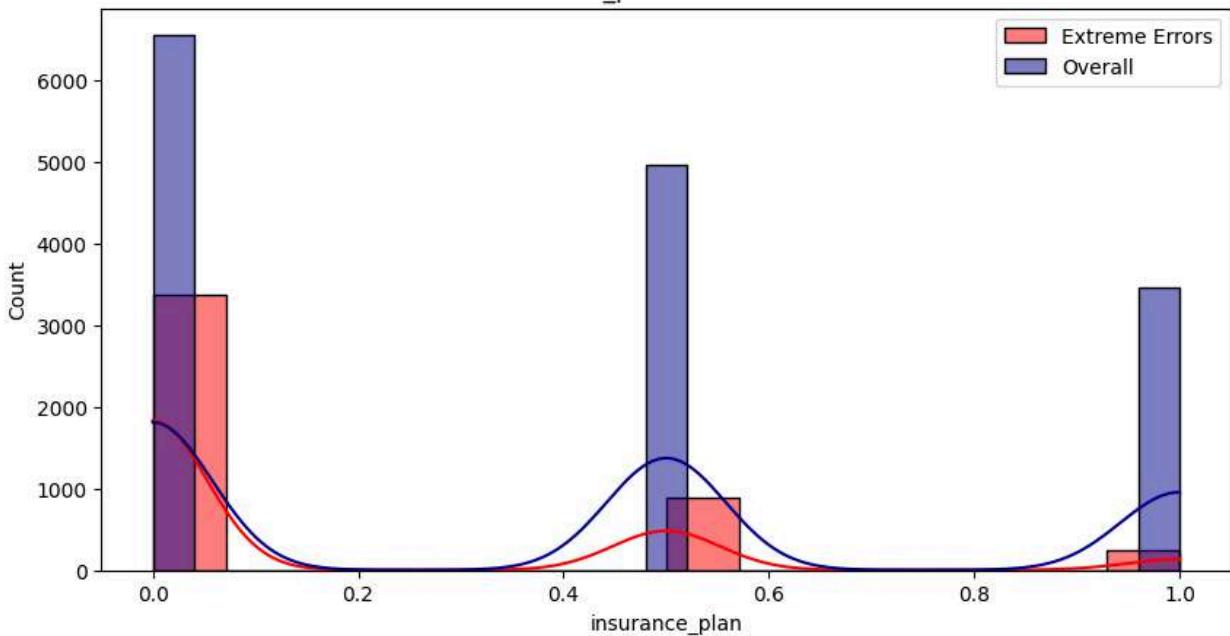
```
plt.title(f'Distribution of {feature} for Extreme Errors vs Overall')
plt.show()
```



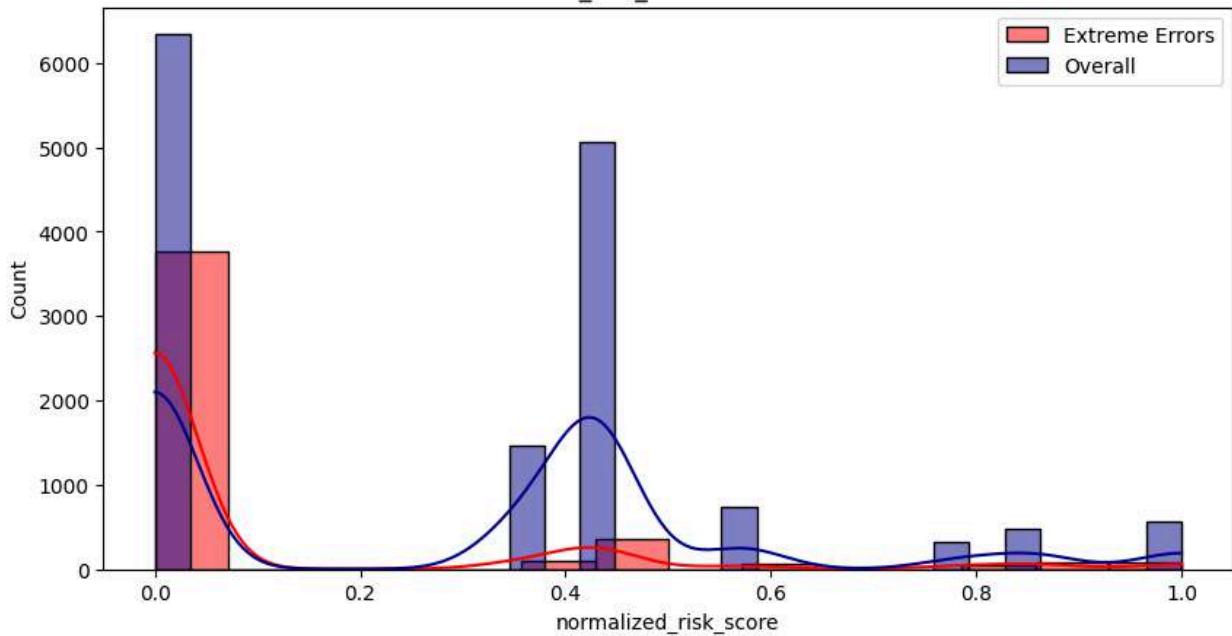
Distribution of income\_lakhs for Extreme Errors vs Overall



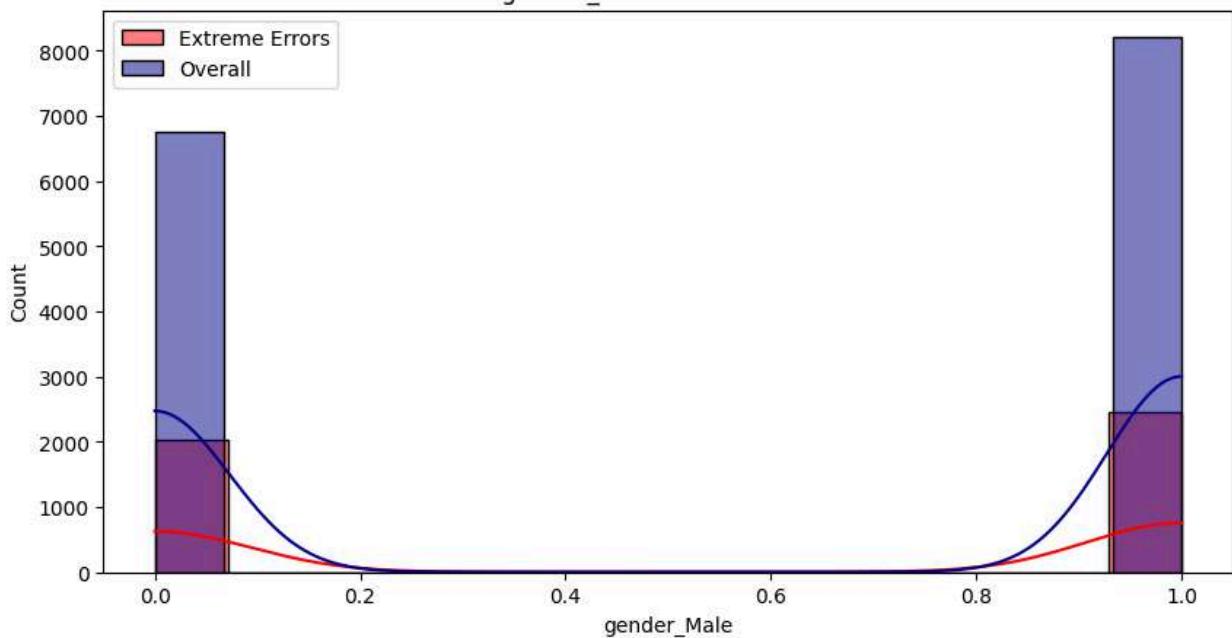
Distribution of insurance\_plan for Extreme Errors vs Overall



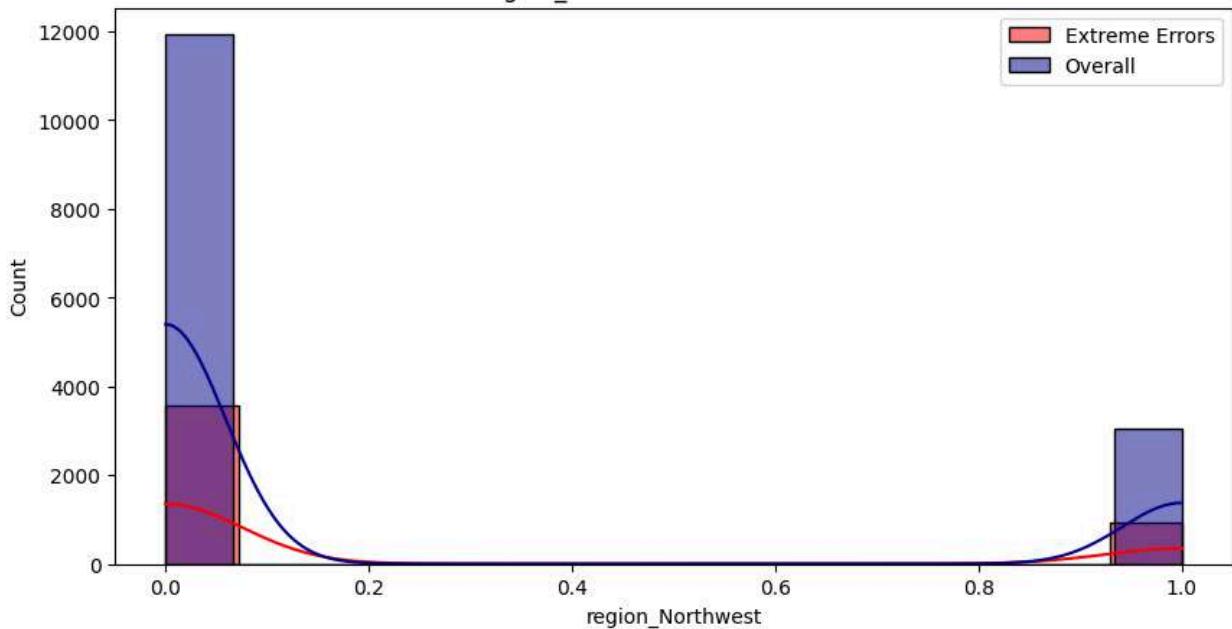
Distribution of normalized\_risk\_score for Extreme Errors vs Overall



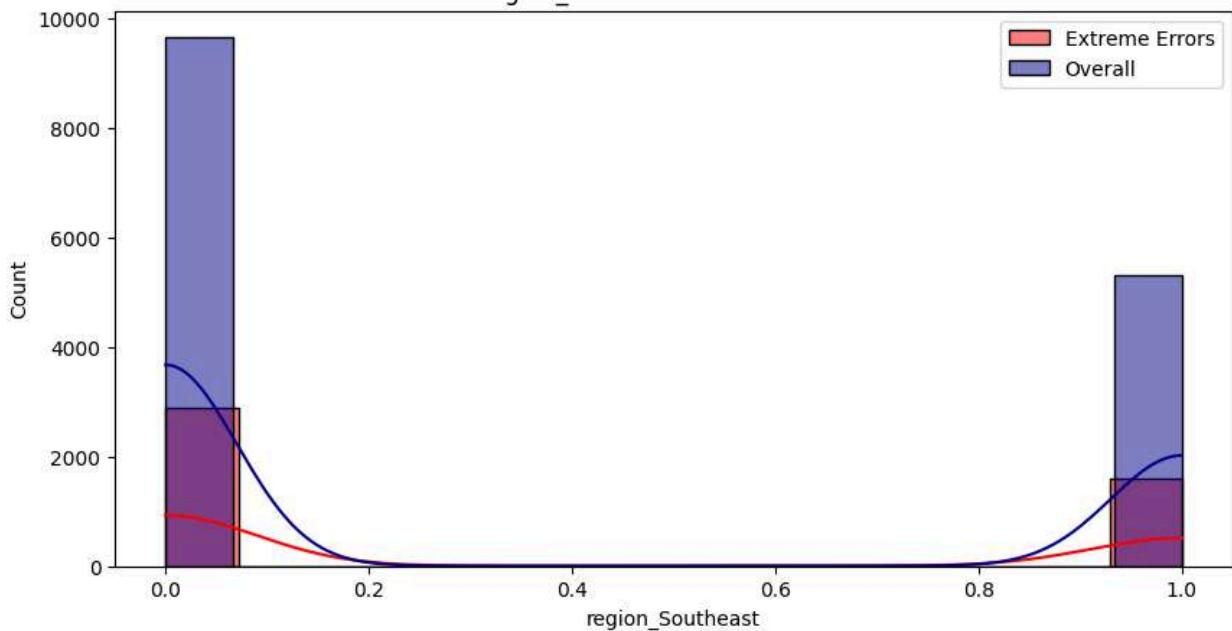
Distribution of gender\_Male for Extreme Errors vs Overall



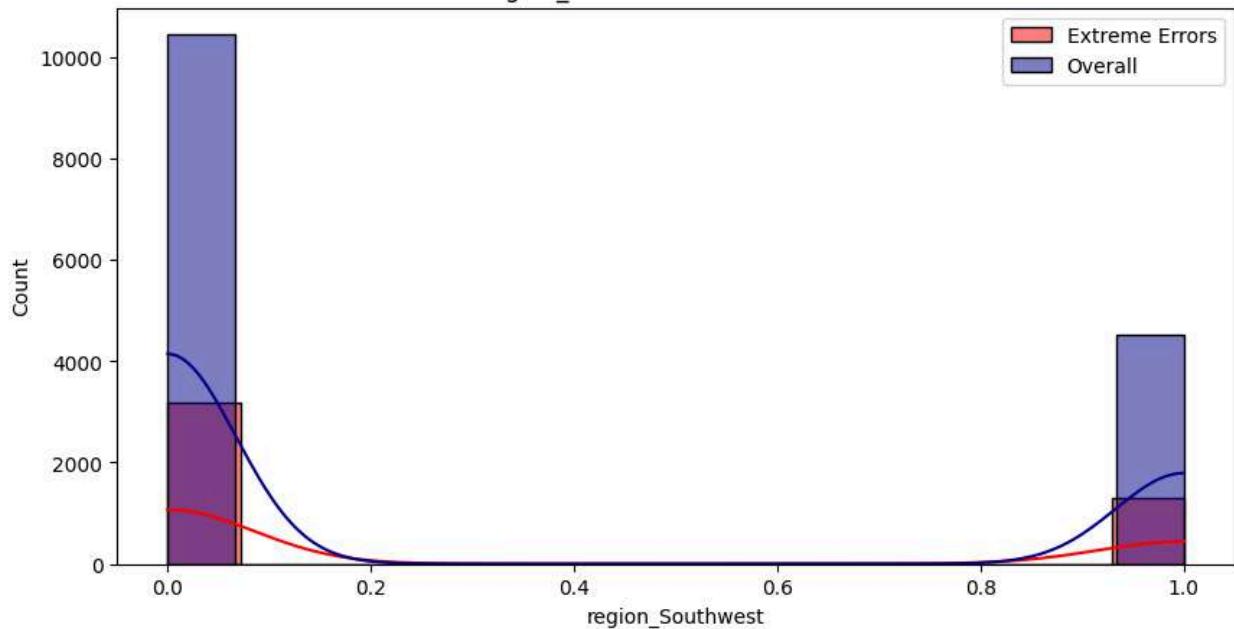
Distribution of region\_Northwest for Extreme Errors vs Overall



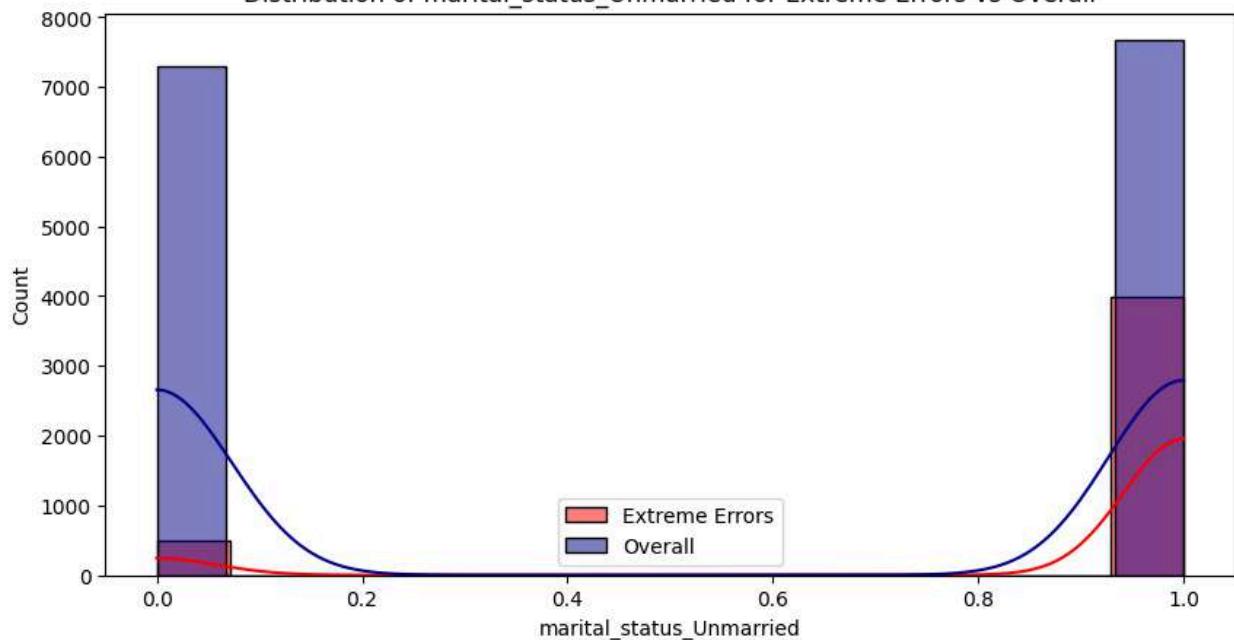
Distribution of region\_Southeast for Extreme Errors vs Overall



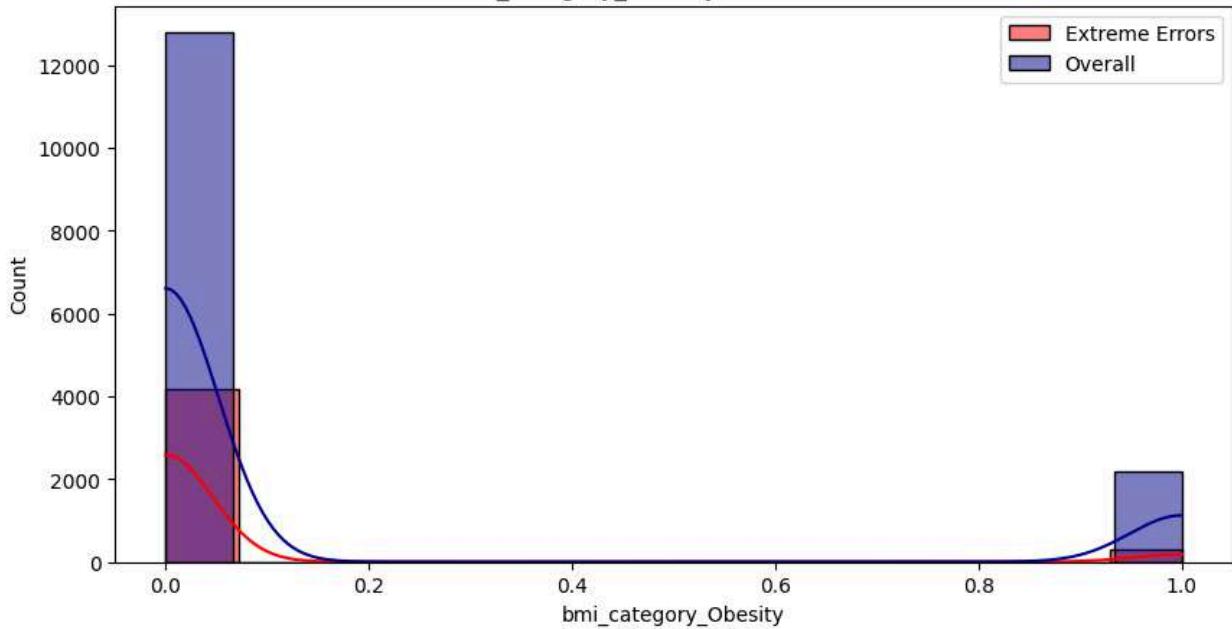
Distribution of region\_Southwest for Extreme Errors vs Overall



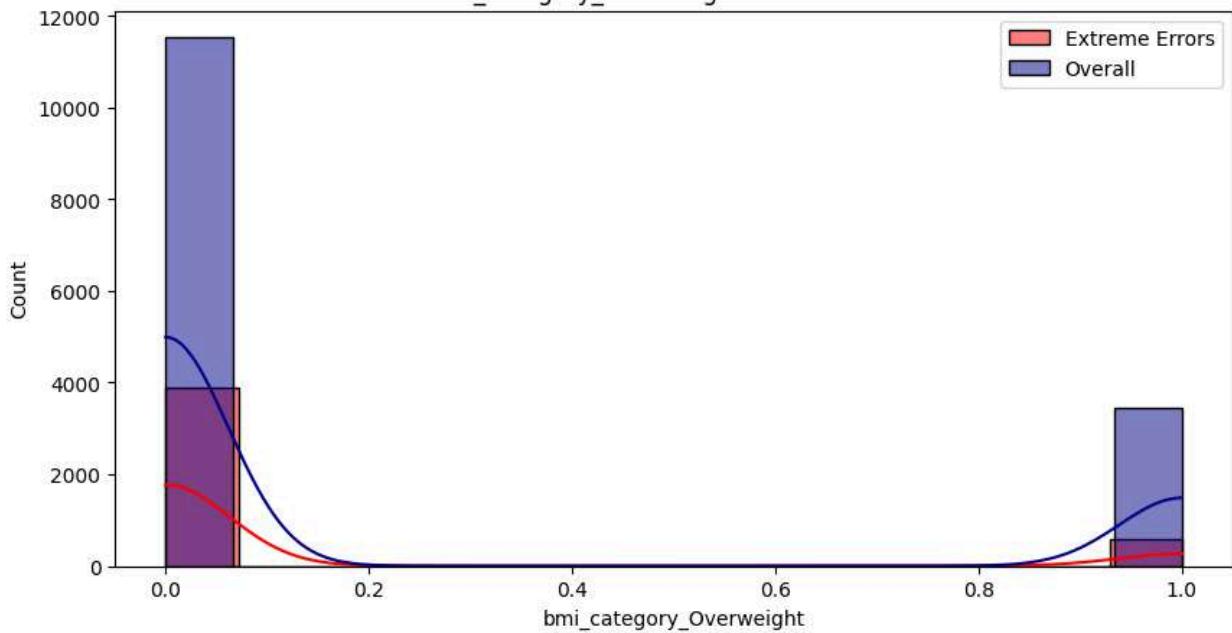
Distribution of marital\_status\_Unmarried for Extreme Errors vs Overall



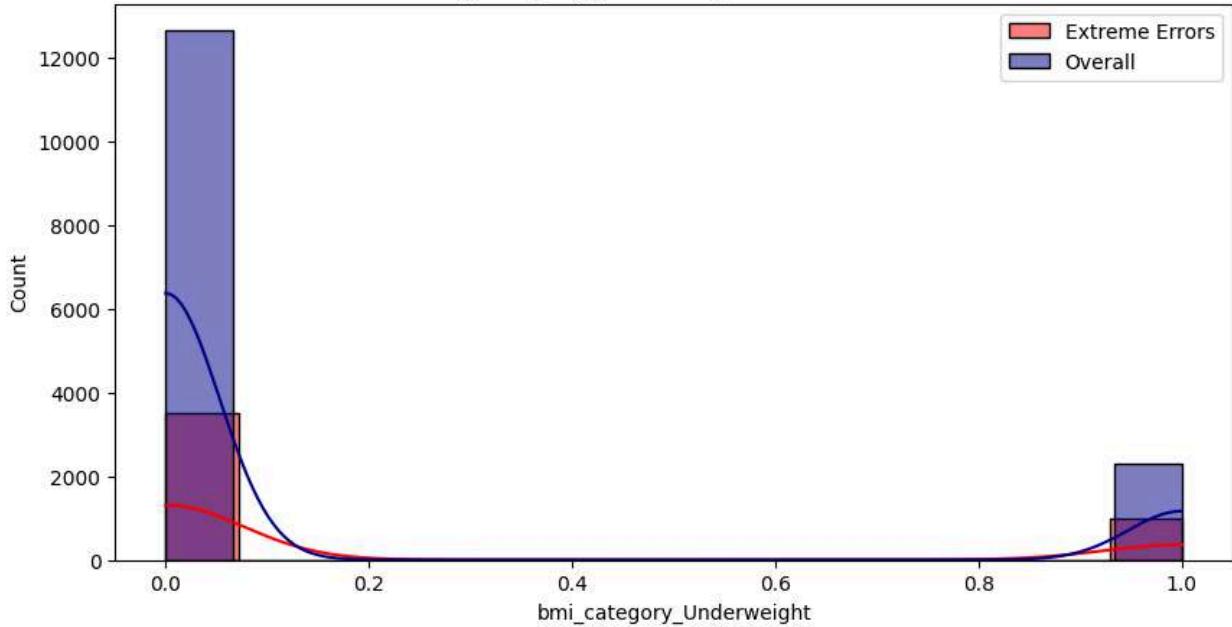
Distribution of bmi\_category\_Obesity for Extreme Errors vs Overall



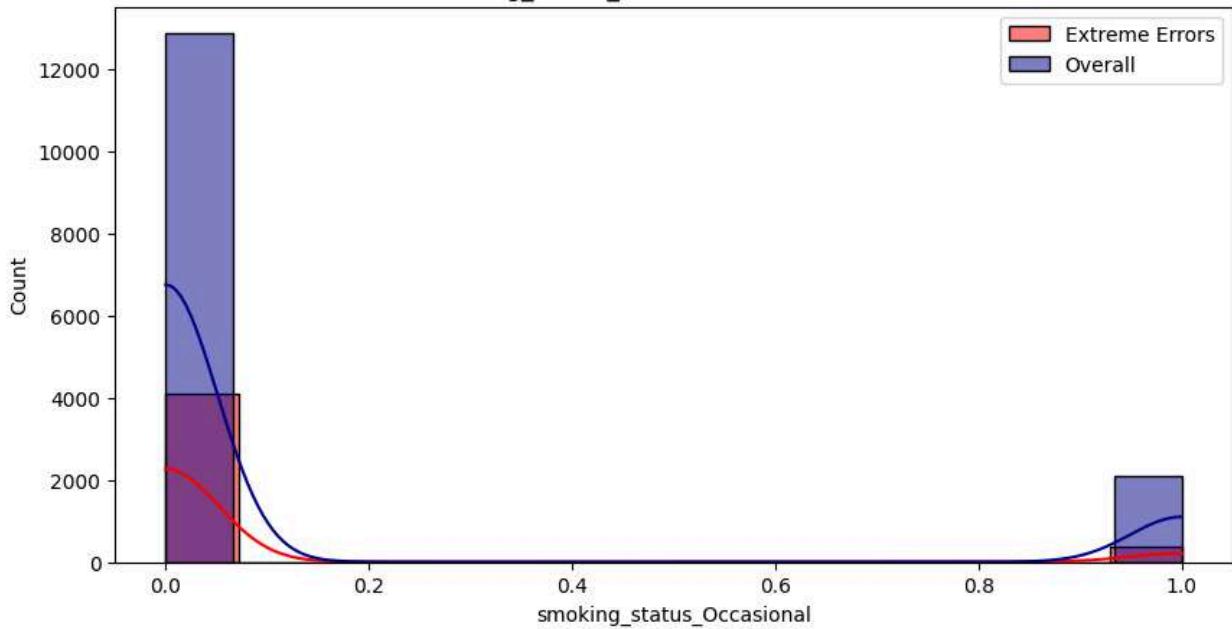
Distribution of bmi\_category\_Overweight for Extreme Errors vs Overall



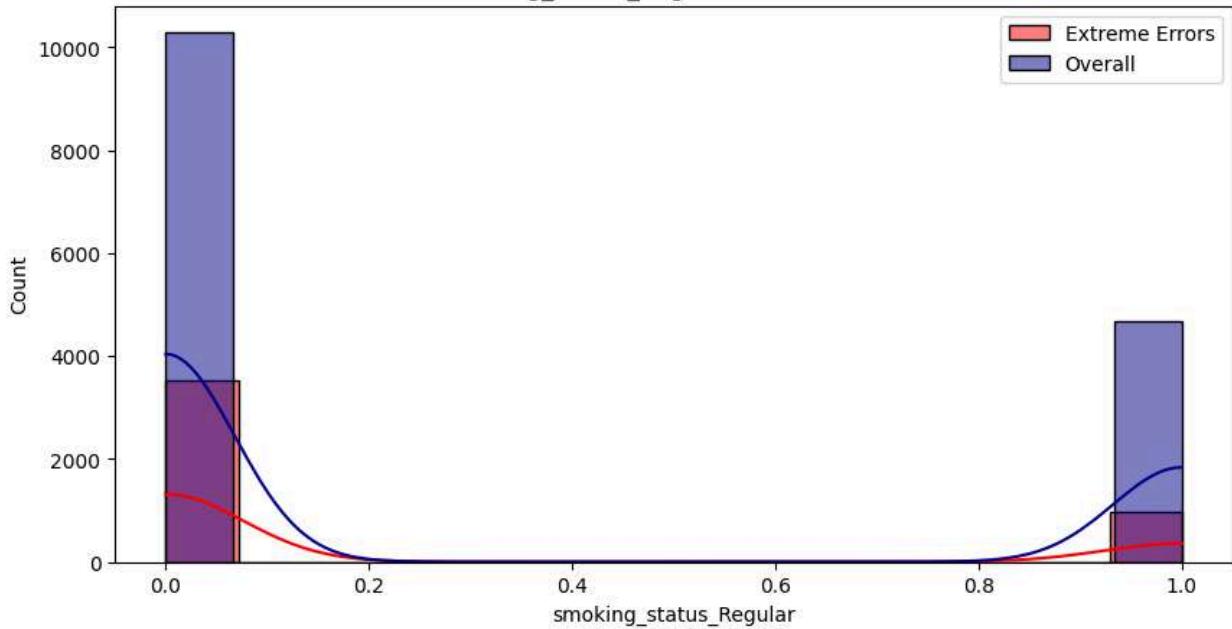
Distribution of bmi\_category\_Underweight for Extreme Errors vs Overall



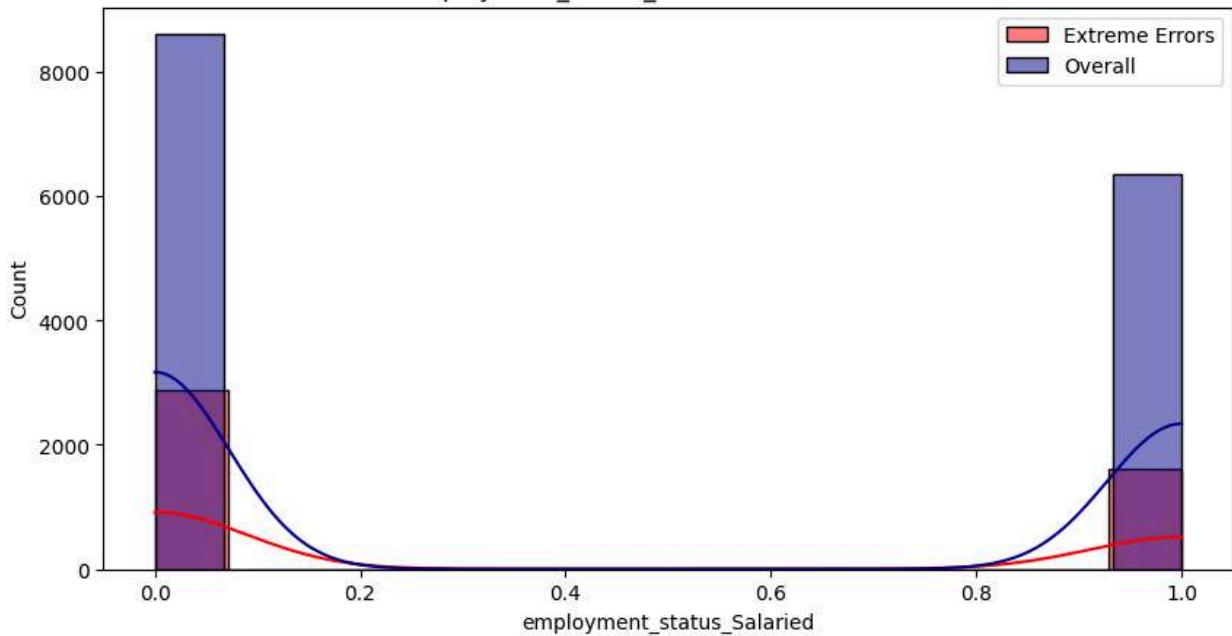
Distribution of smoking\_status\_Occasional for Extreme Errors vs Overall

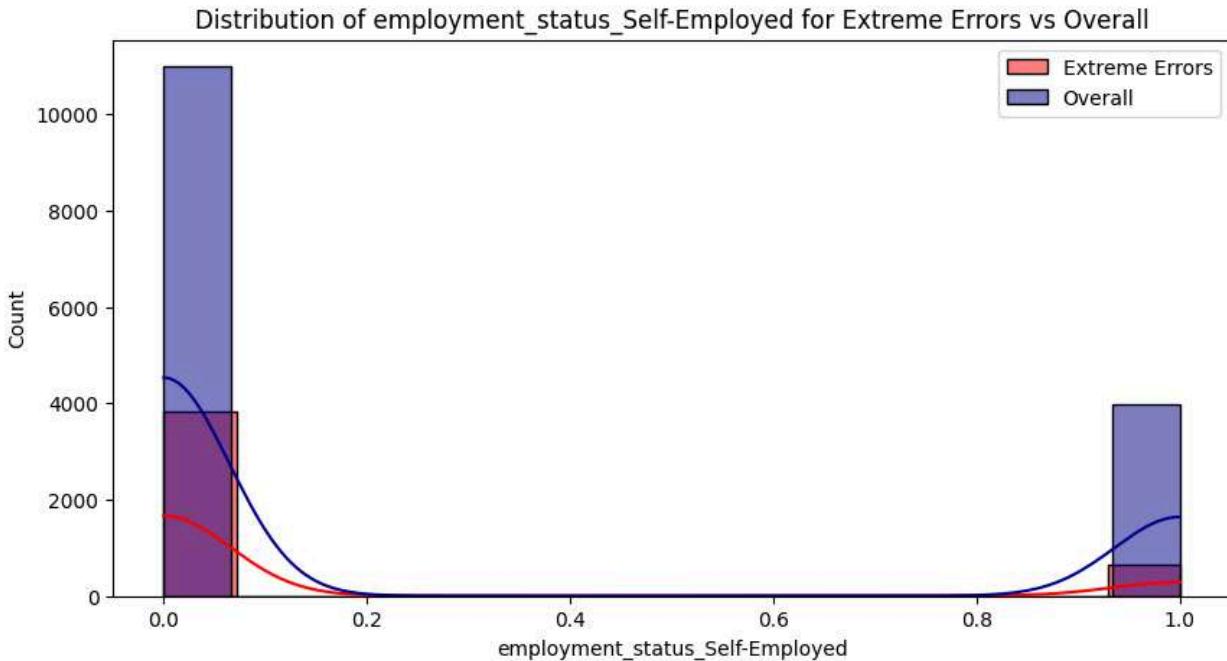


Distribution of smoking\_status-Regular for Extreme Errors vs Overall



Distribution of employment\_status\_Salaried for Extreme Errors vs Overall





Why this analysis is IMPORTANT ??

This tells you:

- Where the model breaks
- Which features cause trouble
- What to fix next

Looking into the above **feature** insights, we can see that the distribution (here it is path of red & blue line) is almost similar in every features except the **age**, and to be specific - **lower range of age**.

## 5G - Reverse Scaling

```
In [80]: extreme_errors_df['income_level']=-1      # since we dropped this column previously
```

```
In [81]: df_reversed = pd.DataFrame()
df_reversed[cols_to_scale] = scaler.inverse_transform(extreme_errors_df[cols_to_scale])
df_reversed.head(10)
```

Out[81]:

	age	number_of_dependants	income_level	income_lakhs	insurance_plan
0	23.0	1.0	-2.0	14.0	1.0
1	19.0	1.0	-2.0	4.0	1.0
2	18.0	1.0	-2.0	3.0	1.0
3	21.0	0.0	-2.0	25.0	1.0
4	24.0	1.0	-2.0	8.0	2.0
5	19.0	0.0	-2.0	7.0	1.0
6	22.0	1.0	-2.0	9.0	2.0
7	20.0	1.0	-2.0	28.0	1.0
8	25.0	0.0	-2.0	8.0	1.0
9	25.0	0.0	-2.0	8.0	1.0

In [82]: `df_reversed.describe()`

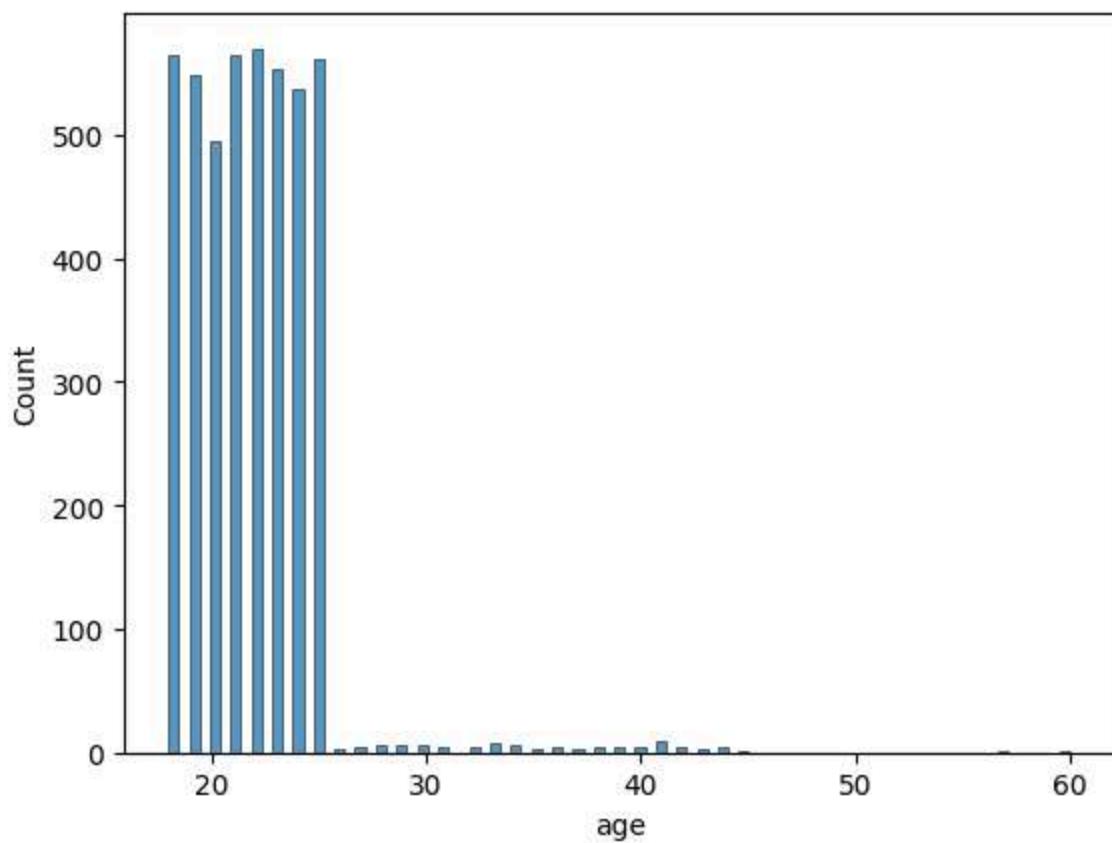
Out[82]:

	age	number_of_dependants	income_level	income_lakhs	insur
<b>count</b>	4487.000000	4487.000000	4487.0	4487.000000	4487.
<b>mean</b>	21.804992	0.739247	-2.0	21.182527	21.182527
<b>std</b>	3.172355	0.968855	0.0	20.598596	20.598596
<b>min</b>	18.000000	0.000000	-2.0	1.000000	1.000000
<b>25%</b>	20.000000	0.000000	-2.0	6.000000	6.000000
<b>50%</b>	22.000000	0.000000	-2.0	15.000000	15.000000
<b>75%</b>	24.000000	1.000000	-2.0	30.000000	30.000000
<b>max</b>	60.000000	5.000000	-2.0	100.000000	100.000000

We can see above that 75% of 'age' is <= 24

In [83]: `sns.histplot(df_reversed.age)`

Out[83]: <Axes: xlabel='age', ylabel='Count'>



```
In [84]: df_reversed['age'].quantile(0.97)
```

```
Out[84]: 25.0
```

This shows that majority (around 97%) of the **extreme errors** are coming from **young age group (i.e. <=25 years of age)**.

We need to maybe build a separate model for this segment.

## 6. Model Segmentation

```
df_youth = df_m[df_m.Age<=25]  
df_mature = df_m[df_m.Age>25]  
  
df_youth.to_excel("prem_youth.xlsx", index = False)  
df_mature.to_excel("prem_mature.xlsx", index = False)
```

## 7. Model Retraining

- Divided the data into 2 parts : **Youth**(age <= 25) and **Mature**(age > 25)
- Added the additional data that displays **Genetic Risk**, for better model training
- Created an artifact folder that has **model & scaler** joblib file for the **youth & mature**

## 8. Building The App

### Why Streamlit is popular ??

- Extremely easy
- Fast to build
- Perfect for ML demos
- Recruiter-friendly
- No web dev skills needed

That's why it's the #1 tool for ML portfolios.

## **What Streamlit is commonly used for ??**

- ML model demos
- Data dashboards
- What-if analysis tools
- Internal business tools

## **9. Deployed The App**

<https://insurance-premium-aiml.streamlit.app/>